

Final Project csc 380

May 8, 2025

1 CSC 380 Final Project

Contributor: Stephin Tomson

1.1 Analyzing the Correlation Between Streetlights and Property-Related Crimes in Tucson

1.1.1 Introduction

Problem Statement

Crime prevention is a critical aspect of urban planning. Adequate street lighting is commonly assumed to deter criminal activity and enhance public safety. However, its actual impact—particularly on property-related offenses—remains underexplored in Tucson.

Hypothesis

Areas with fewer streetlights are more prone to property-related crimes, both during the day and at night.

Objective

This project aims to:

- Determine if a measurable correlation exists between streetlight density and property-related crime rates in Tucson.
 - Use machine learning models to predict crime risk in underlit areas, helping to inform future urban infrastructure planning.
-

1.1.2 Data Gathering and Cleaning

Overview

Data for this project was sourced from public Tucson datasets in CSV format and includes: - Arrest records (Tucson Police Department) - Streetlight infrastructure - Neighborhood income data - Parcel-level land/property information

Each dataset underwent the following cleaning and preprocessing steps:

- **Date & Time Parsing**
Transformed date and time strings into datetime objects to enable accurate filtering and grouping (e.g., by hour or time of day).
- **Essential Column Selection**
Unnecessary or irrelevant columns were dropped to focus on key attributes such as location, demographic info, and charge descriptions.
- **Column Renaming**
Renamed columns for clarity and consistency (e.g., $X \rightarrow *_location_x$), ensuring uniform naming conventions across datasets.
- **Time Period Classification**
Derived a `time_period` column (Morning, Afternoon, Evening, Night) using the hour extracted from arrest time.
- **Filtering Active Records**
Kept only “Active” entries from the streetlight dataset, and removed rows with critical null values in other datasets.
- **Grid-Based Spatial Indexing**
Mapped coordinates to spatial grid cells (approx. 1km²) to facilitate neighborhood-level aggregation and merging across all datasets using `grid_cell_id`.

These steps ensured that each dataset was well-structured, consistent, and ready for integration and analysis.

```
[2]: !pip install s3fs
```

```
Requirement already satisfied: s3fs in /usr/local/lib/python3.11/dist-packages
(2025.3.2)
Requirement already satisfied: aiobotocore<3.0.0,>=2.5.4 in
/usr/local/lib/python3.11/dist-packages (from s3fs) (2.22.0)
Requirement already satisfied: fsspec==2025.3.2.* in
/usr/local/lib/python3.11/dist-packages (from s3fs) (2025.3.2)
Requirement already satisfied: aiohttp!=4.0.0a0,!4.0.0a1 in
/usr/local/lib/python3.11/dist-packages (from s3fs) (3.11.15)
Requirement already satisfied: aioitertools<1.0.0,>=0.5.1 in
/usr/local/lib/python3.11/dist-packages (from aiobotocore<3.0.0,>=2.5.4->s3fs)
(0.12.0)
Requirement already satisfied: botocore<1.37.4,>=1.37.2 in
/usr/local/lib/python3.11/dist-packages (from aiobotocore<3.0.0,>=2.5.4->s3fs)
(1.37.3)
Requirement already satisfied: python-dateutil<3.0.0,>=2.1 in
/usr/local/lib/python3.11/dist-packages (from aiobotocore<3.0.0,>=2.5.4->s3fs)
(2.9.0.post0)
Requirement already satisfied: jmespath<2.0.0,>=0.7.1 in
/usr/local/lib/python3.11/dist-packages (from aiobotocore<3.0.0,>=2.5.4->s3fs)
(1.0.1)
```

Requirement already satisfied: multidict<7.0.0,>=6.0.0 in /usr/local/lib/python3.11/dist-packages (from aiobotocore<3.0.0,>=2.5.4->s3fs) (6.4.3)

Requirement already satisfied: wrapt<2.0.0,>=1.10.10 in /usr/local/lib/python3.11/dist-packages (from aiobotocore<3.0.0,>=2.5.4->s3fs) (1.17.2)

Requirement already satisfied: aiohappyeyeballs>=2.3.0 in /usr/local/lib/python3.11/dist-packages (from aiohttp!=4.0.0a0,!4.0.0a1->s3fs) (2.6.1)

Requirement already satisfied: aiosignal>=1.1.2 in /usr/local/lib/python3.11/dist-packages (from aiohttp!=4.0.0a0,!4.0.0a1->s3fs) (1.3.2)

Requirement already satisfied: attrs>=17.3.0 in /usr/local/lib/python3.11/dist-packages (from aiohttp!=4.0.0a0,!4.0.0a1->s3fs) (25.3.0)

Requirement already satisfied: frozenlist>=1.1.1 in /usr/local/lib/python3.11/dist-packages (from aiohttp!=4.0.0a0,!4.0.0a1->s3fs) (1.6.0)

Requirement already satisfied: propcache>=0.2.0 in /usr/local/lib/python3.11/dist-packages (from aiohttp!=4.0.0a0,!4.0.0a1->s3fs) (0.3.1)

Requirement already satisfied: yarl<2.0,>=1.17.0 in /usr/local/lib/python3.11/dist-packages (from aiohttp!=4.0.0a0,!4.0.0a1->s3fs) (1.20.0)

Requirement already satisfied: urllib3!=2.2.0,<3,>=1.25.4 in /usr/local/lib/python3.11/dist-packages (from botocore<1.37.4,>=1.37.2->aiobotocore<3.0.0,>=2.5.4->s3fs) (2.4.0)

Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.11/dist-packages (from python-dateutil<3.0.0,>=2.1->aiobotocore<3.0.0,>=2.5.4->s3fs) (1.17.0)

Requirement already satisfied: idna>=2.0 in /usr/local/lib/python3.11/dist-packages (from yarl<2.0,>=1.17.0->aiohttp!=4.0.0a0,!4.0.0a1->s3fs) (3.10)

```
[3]: # Importing all of the necessary packages that are going to be utilized,
      ↪ throughout the Notebook.
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import mean_squared_error, classification_report
import folium
from datetime import datetime
import geopandas as gpd
from scipy import stats
```

```
import s3fs
from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

###Importing data

```
[4]: tuc_arrest_data = pd.read_csv('/content/drive/MyDrive/csc380_data/
↳Tucson_Police_Arrests_-_2021_-_Open_Data.csv', low_memory=False,
↳encoding=None)
tuc_light_data = pd.read_csv('/content/drive/MyDrive/csc380_data/
↳Streetlights_-_City_of_Tucson_-_Open_Data.csv', low_memory=False,
↳encoding=None)
tuc_neighbourhood_income = pd.read_csv('/content/drive/MyDrive/csc380_data/
↳Neighborhood_Income.csv', low_memory=False, encoding=None)
tuc_parcel_data = pd.read_csv('/content/drive/MyDrive/csc380_data/
↳Parcels_-_Regional.csv', low_memory=False, encoding=None)
```

1.1.3 Printing Data Types

```
[5]: print("Arrests Data Info:") # Print a header for the arrests data information
print(tuc_arrest_data.info()) # Print the information summary of the arrests
↳DataFrame
print("\nArrests Data Head:") # Print a header for the first few rows of the
↳arrests data
print(tuc_arrest_data.head()) # Print the first 5 rows of the arrests DataFrame
```

Arrests Data Info:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 45421 entries, 0 to 45420

Data columns (total 39 columns):

#	Column	Non-Null Count	Dtype
0	OBJECTID	45421 non-null	int64
1	X	45421 non-null	float64
2	Y	45421 non-null	float64
3	arre_id	45421 non-null	int64
4	case_id	45421 non-null	int64
5	agency	45421 non-null	object
6	date_arr	45421 non-null	object
7	time_arr	45421 non-null	int64
8	datetime_arr	45421 non-null	object
9	MONTH_ARR	45421 non-null	object
10	YEAR_ARR	45421 non-null	int64
11	DOW_ARR	45421 non-null	object

12	TIME_ARRST	45421	non-null	int64
13	age	45421	non-null	object
14	race	45421	non-null	object
15	sex	45421	non-null	object
16	ethnicity	45421	non-null	object
17	arr_type	45421	non-null	object
18	neighborhd	45421	non-null	object
19	ADDRESS_PUBLIC	45416	non-null	object
20	city	43901	non-null	object
21	state	43899	non-null	object
22	zip	43901	non-null	object
23	arr_chrg	45421	non-null	object
24	chrgdesc	45421	non-null	object
25	chrg_cnt	45421	non-null	int64
26	fel_misd	45421	non-null	object
27	chrg_seq	45421	non-null	int64
28	APPSTATE	45421	non-null	int64
29	LOC_STATUS	45421	non-null	object
30	WARD	44591	non-null	float64
31	NHA_NAME	32086	non-null	object
32	TMSECT	0	non-null	float64
33	DIVISION	45065	non-null	object
34	DIVISION_NO	44576	non-null	object
35	DIVSECT	44576	non-null	object
36	TRSQ	45049	non-null	object
37	City_geo	44995	non-null	object
38	ADDRESS_100BLK	45416	non-null	object

dtypes: float64(4), int64(9), object(26)

memory usage: 13.5+ MB

None

Arrests Data Head:

	OBJECTID	X	Y	arre_id	case_id	agency	\
0	1	9.900089e+05	470751.276735	2021000107	2101020104	TPD	
1	2	9.900089e+05	470751.276735	2021000107	2101020104	TPD	
2	3	9.900089e+05	470751.276735	2021000107	2101020104	TPD	
3	4	1.053154e+06	443419.380064	2021000110	2101020138	TPD	
4	5	1.053154e+06	443419.380064	2021000110	2101020138	TPD	

		date_arr	time_arr		datetime_arr	MONTH_ARR	...	\
0	2021/01/02	00:00:00+00	1731	2021/01/02	17:31:00+00	01-Jan	...	
1	2021/01/02	00:00:00+00	1731	2021/01/02	17:31:00+00	01-Jan	...	
2	2021/01/02	00:00:00+00	1731	2021/01/02	17:31:00+00	01-Jan	...	
3	2021/01/02	00:00:00+00	1844	2021/01/02	18:44:00+00	01-Jan	...	
4	2021/01/02	00:00:00+00	1844	2021/01/02	18:44:00+00	01-Jan	...	

	LOC_STATUS	WARD	NHA_NAME	TMSECT		DIVISION	DIVISION_NO	\
0	GEOCODED	3.0	NaN	NaN	Operations Division West		T2	

1	GEOCODED	3.0	NaN	NaN	Operations Division West	T2
2	GEOCODED	3.0	NaN	NaN	Operations Division West	T2
3	GEOCODED	2.0	Eastside	NaN	Operations Division East	T4
4	GEOCODED	2.0	Eastside	NaN	Operations Division East	T4

	DIVSECT	TRSQ	City_geo	ADDRESS_100BLK
0	T203	13S13E24NW	TUCSON	4598 N ORACLE RD
1	T203	13S13E24NW	TUCSON	4598 N ORACLE RD
2	T203	13S13E24NW	TUCSON	4598 N ORACLE RD
3	T406	14S15E14NE	TUCSON	10198 E ESSEX VILLAGE DR
4	T406	14S15E14NE	TUCSON	10198 E ESSEX VILLAGE DR

[5 rows x 39 columns]

```
[6]: print("Lights Data Info:") # Print a header for the streetlights data
      ↪ information
      print(tuc_light_data.info()) # Print the information summary of the
      ↪ streetlights DataFrame
      print("\nLights Data Head:") # Print a header for the first few rows of the
      ↪ streetlights data
      print(tuc_light_data.head()) # Print the first 5 rows of the streetlights
      ↪ DataFrame
```

Lights Data Info:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 22781 entries, 0 to 22780

Data columns (total 21 columns):

#	Column	Non-Null Count	Dtype
---	-----	-----	-----
0	X	22768 non-null	float64
1	Y	22768 non-null	float64
2	OBJECTID	22781 non-null	int64
3	Model	22423 non-null	object
4	Type	22446 non-null	object
5	Bulb_Type	22455 non-null	object
6	Wattage	22338 non-null	float64
7	Voltage	22235 non-null	float64
8	Address_Number	18468 non-null	float64
9	Street	19451 non-null	object
10	City	22542 non-null	object
11	Light_Fixture_Theme	22781 non-null	object
12	DATASOURCE	22781 non-null	object
13	Permit_Number	562 non-null	object
14	CartegraphID	22768 non-null	object
15	SHAPE	0 non-null	float64
16	Retired	8 non-null	object
17	Status	22768 non-null	object

```

18 MacID 19638 non-null object
19 TEP_Account_Number 22183 non-null float64
20 Power_Pedestal_ID 22191 non-null float64
dtypes: float64(8), int64(1), object(12)
memory usage: 3.7+ MB
None

```

Lights Data Head:

	X	Y	OBJECTID	Model	Type	Bulb_Type	\
0	1.001233e+06	421018.579396	1	ATBM D R3	Autobahn	LED	
1	1.001142e+06	420902.066601	2	ATBM D R3	Autobahn	LED	
2	1.001234e+06	420785.376969	3	ATBM D R3	Autobahn	LED	
3	1.001143e+06	420667.403543	4	ATBM D R3	Autobahn	LED	
4	1.001237e+06	420582.028543	5	ATBM D R3	Autobahn	LED	

	Wattage	Voltage	Address_Number	Street	...	Light_Fixture_Theme	\
0	95.0	480.0	5425.0	S Campbell Av	...	Other	
1	95.0	480.0	5434.0	S Campbell Av	...	Other	
2	95.0	480.0	5441.0	S Campbell Av	...	Other	
3	95.0	480.0	5454.0	S Campbell Av	...	Other	
4	95.0	480.0	5457.0	S Campbell Av	...	Other	

	DATASOURCE	Permit_Number	CartegraphID	SHAPE	Retired	Status	\
0	TDOT_STREETLIGHTS	NaN	51	NaN	NaN	Active	
1	TDOT_STREETLIGHTS	NaN	52	NaN	NaN	Active	
2	TDOT_STREETLIGHTS	NaN	53	NaN	NaN	Active	
3	TDOT_STREETLIGHTS	NaN	54	NaN	NaN	Active	
4	TDOT_STREETLIGHTS	NaN	55	NaN	NaN	Active	

	MacID	TEP_Account_Number	Power_Pedestal_ID
0	00F14C41	3.710003e+09	514.0
1	00F10A57	3.710003e+09	514.0
2	00F16079	3.710003e+09	514.0
3	00F10902	3.710003e+09	514.0
4	00F15E6B	3.710003e+09	514.0

[5 rows x 21 columns]

```

[7]: print("Income Data Info:") # Print a header for the neighborhood income data_
      ↪information
      print(tuc_neighbourhood_income.info()) # Print the information summary of the_
      ↪neighborhood income DataFrame
      print("\nIncome Data Head:") # Print a header for the first few rows of the_
      ↪neighborhood income data
      print(tuc_neighbourhood_income.head()) # Print the first 5 rows of the_
      ↪neighborhood income DataFrame

```

Income Data Info:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 159 entries, 0 to 158

Columns: 168 entries, OBJECTID to WLTHINDXCY

dtypes: float64(5), int64(157), object(6)

memory usage: 208.8+ KB

None

Income Data Head:

	OBJECTID	NAME	WARD	DATASOURCE	ID	sourceCountry	\
0	1	A Mountain	1	NEIGHBORHOODS	0	US	
1	2	Adelanto	3	NEIGHBORHOODS	1	US	
2	3	Alvernon Heights	5	NEIGHBORHOODS	2	US	
3	4	Amphi	3	NEIGHBORHOODS	3	US	
4	5	Armory Park	6	NEIGHBORHOODS	4	US	

	ENRICH_FID	aggregationMethod	\
0	1	BlockApportionment:US.BlockGroups	
1	2	BlockApportionment:US.BlockGroups	
2	3	BlockApportionment:US.BlockGroups	
3	4	BlockApportionment:US.BlockGroups	
4	5	BlockApportionment:US.BlockGroups	

	populationToPolygonSizeRating	apportionmentConfidence	...	AGGDIA75CY	\
0	2.191	2.576	...	1590160	
1	2.191	2.576	...	154598	
2	2.191	2.576	...	172634	
3	2.191	2.576	...	2760918	
4	2.191	2.576	...	3785750	

	ID_1	sourceCountry_1	ENRICH_FID_1	aggregationMethod_1	\
0	0	US	1	BlockApportionment:US.BlockGroups	
1	1	US	2	BlockApportionment:US.BlockGroups	
2	2	US	3	BlockApportionment:US.BlockGroups	
3	3	US	4	BlockApportionment:US.BlockGroups	
4	4	US	5	BlockApportionment:US.BlockGroups	

	populationToPolygonSizeRating_1	apportionmentConfidence_1	HasData_1	\
0	2.191	2.576	1	
1	2.191	2.576	1	
2	2.191	2.576	1	
3	2.191	2.576	1	
4	2.191	2.576	1	

	TOTHH_CY	WLTHINDXCY
0	1103	32
1	117	28
2	99	26


```
3      3105      20
4      1223      48
```

[5 rows x 168 columns]

```
[8]: print("Parcel Data Info:") # Print a header for the parcel data information
      print(tuc_parcel_data.info()) # Print the information summary of the parcel_
      ↪ DataFrame
      print("\nParcel Data Head:") # Print a header for the first few rows of the_
      ↪ parcel data
      tuc_parcel_data.head() # Display the first 5 rows of the parcel DataFrame
```

Parcel Data Info:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 443275 entries, 0 to 443274

Data columns (total 50 columns):

#	Column	Non-Null Count	Dtype
---	-----	-----	-----
0	OBJECTID	443275 non-null	int64
1	PARCEL	443275 non-null	object
2	GISAREA	443275 non-null	float64
3	GISACRES	443275 non-null	float64
4	X_HPGN	443275 non-null	float64
5	Y_HPGN	443275 non-null	float64
6	LON	443275 non-null	float64
7	LAT	443275 non-null	float64
8	LOT_R	383138 non-null	object
9	LINK	443275 non-null	object
10	TRS_OL	443275 non-null	object
11	MP_OL	353970 non-null	float64
12	SEQ_NUM_S	388490 non-null	float64
13	JURIS_OL	443275 non-null	object
14	CURZONE_OL	443258 non-null	object
15	ADDRESS_OL	386148 non-null	object
16	ADR_STATUS	443275 non-null	object
17	SEQ_NUM_D	439054 non-null	float64
18	PARCEL_USE	441936 non-null	float64
19	LANDMEAS	441936 non-null	float64
20	LANDUNIT	441936 non-null	object
21	LASTCHANGE	441936 non-null	object
22	LEGAL1	441936 non-null	object
23	LEGAL2	141539 non-null	object
24	LEGAL3	21788 non-null	object
25	LEGAL4	6727 non-null	object
26	LEGAL5	11649 non-null	object
27	LOT	383411 non-null	object
28	MAIL1	441936 non-null	object
29	MAIL2	439185 non-null	object

```

30 MAIL3          439092 non-null object
31 MAIL4          88858 non-null object
32 MAIL5          13060 non-null object
33 MP             388545 non-null object
34 PAGE           438282 non-null float64
35 RECORDDATE     434041 non-null float64
36 DOCKET         438282 non-null float64
37 RECTRACT       427636 non-null object
38 SECTMODIF      236605 non-null object
39 TAXAREA        441936 non-null float64
40 ZIP            441827 non-null object
41 ZIP4           441480 non-null object
42 TAXYR          439546 non-null float64
43 LIMNET         439546 non-null float64
44 FCV            439546 non-null float64
45 last_edited_user 443275 non-null object
46 last_edited_date 443275 non-null object
47 PC_RESTRICT    983 non-null object
48 ShapeSTArea    443275 non-null float64
49 ShapeSTLength  443275 non-null float64
dtypes: float64(20), int64(1), object(29)
memory usage: 169.1+ MB
None

```

Parcel Data Head:

```

[8]:  OBJECTID  PARCEL      GISAREA  GISACRES      X_HPGN      Y_HPGN  \
0      1  10101001D  110172.233186  2.529114  979285.912649  487059.242668
1      2  10101002A  114140.158529  2.620201  979275.589076  486317.677724
2      3  10101003E   94614.544440  2.171971  979098.778243  482696.304361
3      4  10101003L   80326.675549  1.843979  978758.273758  482660.756604
4      5  10101003M   79901.424565  1.834217  978437.135645  482651.278309

      LON      LAT  LOT_R      LINK  ...  \
0 -111.012409  32.335766   NaN  HTTPS://GIS.PIMA.GOV/D.HTM?P=10101001D  ...
1 -111.012462  32.333728   NaN  HTTPS://GIS.PIMA.GOV/D.HTM?P=10101002A  ...
2 -111.013134  32.323779   NaN  HTTPS://GIS.PIMA.GOV/D.HTM?P=10101003E  ...
3 -111.014237  32.323689   NaN  HTTPS://GIS.PIMA.GOV/D.HTM?P=10101003L  ...
4 -111.015276  32.323670   NaN  HTTPS://GIS.PIMA.GOV/D.HTM?P=10101003M  ...

      ZIP  ZIP4  TAXYR  LIMNET      FCV  last_edited_user  \
0  00000  0000  2024.0      0.0    3710.0  u142832@CENTRAL
1  85705  1547  2024.0      0.0  1918109.0  u142832@CENTRAL
2  60015  6002  2024.0  436783.0  2964029.0  u142832@CENTRAL
3  85741  3117  2024.0  330439.0  3220744.0      GISPARFAB
4  85715  3808  2024.0  236550.0  1690713.0      GISPARFAB

```

	last_edited_date	PC_RESTRICT	ShapeSTArea	ShapeSTLength
0	2021/01/06 15:20:56+00	NaN	110169.778019	1911.315304
1	2021/01/06 15:20:56+00	NaN	114138.031173	1676.392288
2	2021/01/06 15:20:56+00	NaN	94612.521449	1235.311573
3	2020/10/20 00:14:03+00	NaN	80326.675549	1142.977487
4	2020/10/20 00:49:07+00	NaN	79901.424565	1166.740103

[5 rows x 50 columns]

1.2 Data Cleaning and Pre-Processing

1.2.1 Data Cleaning & Preprocessing Overview

This section cleans and prepares four Tucson datasets for spatial analysis:

- **Arrests:** Parsed dates/times, converted age, extracted arrest hour & time period, renamed key fields.
- **Streetlights:** Filtered for active lights, retained location and device info.
- **Neighborhoods:** Selected income & household data, renamed columns for clarity.
- **Parcels:** Cleaned location and address info, converted record dates.

Each dataset was gridded into 1km² cells.

Aggregated stats (e.g., arrests, lights, parcels) per grid cell.

Merged with ward-level socioeconomic data.

Calculated densities and added cell center coordinates for spatial mapping.

Final dataset is ready for analysis of crime patterns, infrastructure, and inequality.

1.2.2 Extracting data from arrests

This step prepares the arrests dataset by converting types, extracting time features, and selecting relevant columns for analysis.

```
[9]: # Make a copy of the original dataset so we don't change it directly
clean_arrests = tuc_arrest_data.copy()

# Convert the 'age' column to numeric values (invalid ones become NaN)
clean_arrests['age'] = pd.to_numeric(clean_arrests['age'], errors='coerce')

# Convert arrest date to proper datetime format
clean_arrests['arrest_date'] = pd.to_datetime(clean_arrests['date_arr'])

# Copy time and year of arrest into clearer columns
clean_arrests['arrest_time'] = clean_arrests['time_arr']
clean_arrests['arrest_year'] = clean_arrests['YEAR_ARR']

# Extract just the hour from the arrest time (e.g., 1345 → 13)
clean_arrests['arrest_hour'] = clean_arrests['time_arr'].apply(
```

```

        lambda x: int(str(x).zfill(4)[:2]) if pd.notnull(x) else None
    )

    # Categorize the arrest time into parts of the day
    def categorize_time(hour):
        if pd.isna(hour):
            return 'Unknown'
        if 6 <= hour < 12:
            return 'Morning'
        elif 12 <= hour < 17:
            return 'Afternoon'
        elif 17 <= hour < 22:
            return 'Evening'
        return 'Night'

    # Add the new time period column
    clean_arrests['time_period'] = clean_arrests['arrest_hour'].
        ↪ apply(categorize_time)

    # Keep only the columns we care about and rename some for clarity
    essential_columns = [
        'OBJECTID', 'X', 'Y', 'date_arr', 'time_arr', 'YEAR_ARR',
        'age', 'race', 'zip', 'sex', 'ethnicity', 'arr_type',
        'arr_chrg', 'chrgdesc', 'WARD', 'NHA_NAME', 'ADDRESS_100BLK',
        'arrest_date', 'arrest_time', 'arrest_year', 'arrest_hour', 'time_period'
    ]

    clean_arrests = clean_arrests[essential_columns].rename(columns={
        'X': 'arrest_x_cord',
        'Y': 'arrest_y_cord',
        'arr_chrg': 'arrest_charge_code',
        'chrgdesc': 'arrest_charge_description',
        'WARD': 'arrest_ward_number',
        'arr_type': 'arrest_type',
        'ADDRESS_100BLK': 'arrest_block_address',
        'NHA_NAME': 'arrest_neighborhood_name',
        'zip': 'arrest_zip'
    })

    # Show info and a few rows of the cleaned dataset
    print("Arrests Data Info:")
    print(clean_arrests.info())

    print("\nArrests Data Head:")
    print(clean_arrests.head())

```

Arrests Data Info:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 45421 entries, 0 to 45420

Data columns (total 22 columns):

#	Column	Non-Null Count	Dtype
0	OBJECTID	45421 non-null	int64
1	arrest_x_cord	45421 non-null	float64
2	arrest_y_cord	45421 non-null	float64
3	date_arr	45421 non-null	object
4	time_arr	45421 non-null	int64
5	YEAR_ARR	45421 non-null	int64
6	age	44536 non-null	float64
7	race	45421 non-null	object
8	arrest_zip	43901 non-null	object
9	sex	45421 non-null	object
10	ethnicity	45421 non-null	object
11	arrest_type	45421 non-null	object
12	arrest_charge_code	45421 non-null	object
13	arrest_charge_description	45421 non-null	object
14	arrest_ward_number	44591 non-null	float64
15	arrest_neighborhood_name	32086 non-null	object
16	arrest_block_address	45416 non-null	object
17	arrest_date	45421 non-null	datetime64[ns, UTC]
18	arrest_time	45421 non-null	int64
19	arrest_year	45421 non-null	int64
20	arrest_hour	45421 non-null	int64
21	time_period	45421 non-null	object

dtypes: datetime64[ns, UTC](1), float64(4), int64(6), object(11)

memory usage: 7.6+ MB

None

Arrests Data Head:

	OBJECTID	arrest_x_cord	arrest_y_cord	date_arr	time_arr	\
0	1	9.900089e+05	470751.276735	2021/01/02 00:00:00+00	1731	
1	2	9.900089e+05	470751.276735	2021/01/02 00:00:00+00	1731	
2	3	9.900089e+05	470751.276735	2021/01/02 00:00:00+00	1731	
3	4	1.053154e+06	443419.380064	2021/01/02 00:00:00+00	1844	
4	5	1.053154e+06	443419.380064	2021/01/02 00:00:00+00	1844	

	YEAR_ARR	age	race	arrest_zip	sex	...	arrest_charge_code	\
0	2021	28.0	H	85705	M	...	13-3613A	
1	2021	28.0	H	85705	M	...	13-3613A	
2	2021	28.0	H	85705	M	...	13-3613A	
3	2021	15.0	B	85748	M	...	13-1203A2DV	
4	2021	15.0	B	85748	M	...	13-2904A1DV	

arrest_charge_description arrest_ward_number \

0	CONTRIBUTE TO DELINQUENCY/DEPENDENCY OF MINOR ...	3.0
1	CONTRIBUTE TO DELINQUENCY/DEPENDENCY OF MINOR ...	3.0
2	CONTRIBUTE TO DELINQUENCY/DEPENDENCY OF MINOR ...	3.0
3	ASSAULT-CAUSE FEAR OF PHYSICAL INJURY (DOMESTI...	2.0
4	DISORDERLY CONDUCT/DV ...	2.0

	arrest_neighborhood_name	arrest_block_address	\
0	NaN	4598 N ORACLE RD	
1	NaN	4598 N ORACLE RD	
2	NaN	4598 N ORACLE RD	
3	Eastside	10198 E ESSEX VILLAGE DR	
4	Eastside	10198 E ESSEX VILLAGE DR	

	arrest_date	arrest_time	arrest_year	arrest_hour	time_period
0	2021-01-02	00:00:00+00:00	1731	2021	17 Evening
1	2021-01-02	00:00:00+00:00	1731	2021	17 Evening
2	2021-01-02	00:00:00+00:00	1731	2021	17 Evening
3	2021-01-02	00:00:00+00:00	1844	2021	18 Evening
4	2021-01-02	00:00:00+00:00	1844	2021	18 Evening

[5 rows x 22 columns]

1.2.3 Streetlights Dataset Cleaning

This step filters only active streetlights and selects relevant location and device information for further analysis.

```
[10]: # Clean and filter the streetlights dataset

try:
    # Make a copy of the original dataset
    cleanLights = tuc_light_data.copy()

    # Keep only the rows where the streetlight is active
    cleanLights = cleanLights[cleanLights['Status'] == 'Active']

    # Choose the important columns we want to keep
    essential_columns = [
        'X', 'Y',
        'Address_Number',
        'Street',
        'Status',
        'MacID'
    ]

    # Rename columns to more descriptive names
    cleanLights = cleanLights[essential_columns].rename(columns={
        'X': 'light_location_x',
```

```

        'Y': 'light_location_y',
        'Address_Number': 'light_address_number',
        'Street': 'light_street_name',
        'Status': 'light_operational_status',
        'MacID': 'light_device_id'
    })

except Exception as e:
    # Show the error if something goes wrong
    print(f"An error occurred: {str(e)}")
    import traceback
    traceback.print_exc()
# Show info and a few rows of the cleaned dataset
print("Arrests Data Info:")
print(clean_lights.info())

print("\nArrests Data Head:")
print(clean_lights.head())

```

Arrests Data Info:

<class 'pandas.core.frame.DataFrame'>

Index: 22452 entries, 0 to 22780

Data columns (total 6 columns):

#	Column	Non-Null Count	Dtype
0	light_location_x	22452 non-null	float64
1	light_location_y	22452 non-null	float64
2	light_address_number	18265 non-null	float64
3	light_street_name	19221 non-null	object
4	light_operational_status	22452 non-null	object
5	light_device_id	19510 non-null	object

dtypes: float64(3), object(3)

memory usage: 1.2+ MB

None

Arrests Data Head:

	light_location_x	light_location_y	light_address_number	light_street_name	\
0	1.001233e+06	421018.579396	5425.0	S Campbell Av	
1	1.001142e+06	420902.066601	5434.0	S Campbell Av	
2	1.001234e+06	420785.376969	5441.0	S Campbell Av	
3	1.001143e+06	420667.403543	5454.0	S Campbell Av	
4	1.001237e+06	420582.028543	5457.0	S Campbell Av	

	light_operational_status	light_device_id
0	Active	00F14C41
1	Active	00F10A57
2	Active	00F16079

3	Active	00F10902
4	Active	00F15E6B

1.2.4 Neighborhood Income Data Cleaning

This step extracts key socioeconomic indicators from the neighborhood dataset and renames columns for clarity.

```
[11]: # Clean and simplify the neighborhood income dataset

try:
    # Make a copy of the original data to keep it unchanged
    clean_neighborhoods = tuc_neighbourhood_income.copy()

    # Keep only the important columns with economic and location data
    essential_columns = [
        'NAME',
        'WARD',
        'MEDHINC_CY',
        'AVGHINC_CY',
        'PCI_CY',
        'TOTHH_CY',
        'WLTHINDXCY'
    ]

    # Rename the columns to more readable names
    clean_neighborhoods = clean_neighborhoods[essential_columns].
    ↪rename(columns={
        'NAME': 'neigh_full_name',
        'WARD': 'neigh_ward_number',
        'MEDHINC_CY': 'neigh_median_household_income',
        'AVGHINC_CY': 'neigh_average_household_income',
        'PCI_CY': 'neigh_per_capita_income',
        'TOTHH_CY': 'neigh_total_households',
        'WLTHINDXCY': 'neigh_wealth_index'
    })

except Exception as e:
    # Show error message if anything fails
    print(f"An error occurred: {str(e)}")
    import traceback
    traceback.print_exc()

# Show info and a few rows of the cleaned dataset
print("Arrests Data Info:")
print(clean_neighborhoods.info())

print("\nArrests Data Head:")
```



```
print(clean_neighborhoods.head())
```

Arrests Data Info:

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 159 entries, 0 to 158

Data columns (total 7 columns):

#	Column	Non-Null Count	Dtype
0	neigh_full_name	159 non-null	object
1	neigh_ward_number	159 non-null	int64
2	neigh_median_household_income	159 non-null	int64
3	neigh_average_household_income	159 non-null	int64
4	neigh_per_capita_income	159 non-null	int64
5	neigh_total_households	159 non-null	int64
6	neigh_wealth_index	159 non-null	int64

dtypes: int64(6), object(1)

memory usage: 8.8+ KB

None

Arrests Data Head:

	neigh_full_name	neigh_ward_number	neigh_median_household_income \
0	A Mountain	1	39293
1	Adelanto	3	33635
2	Alvernon Heights	5	29762
3	Amphi	3	20213
4	Armory Park	6	36870

	neigh_average_household_income	neigh_per_capita_income \
0	47471	15189
1	41101	11776
2	38963	12634
3	31338	13040
4	62259	37424

	neigh_total_households	neigh_wealth_index
0	1103	32
1	117	28
2	99	26
3	3105	20
4	1223	48

1.2.5 Parcel Data Cleaning

This step extracts relevant location and address information from the parcel dataset and ensures date values are properly formatted.

```
[12]: # Clean and prepare the parcel dataset

try:
    # Make a copy of the original data to avoid changing it
    clean_parcel_data = tuc_parcel_data.copy()

    # Keep only the important columns with location, address, and date info
    essential_columns = [
        'X_HPGN',
        'Y_HPGN',
        'ADDRESS_OL',
        'MAIL2',
        'RECORDDATE',
        'ZIP'
    ]

    # Rename columns to be easier to understand
    clean_parcel_data = clean_parcel_data[essential_columns].rename(columns={
        'X_HPGN': 'parcel_location_x',
        'Y_HPGN': 'parcel_location_y',
        'ADDRESS_OL': 'parcel_official_address',
        'MAIL2': 'parcel_mailing_address',
        'RECORDDATE': 'parcel_record_date',
        'ZIP': 'parcel_zipcode'
    })

    # Convert record date to datetime format
    clean_parcel_data['parcel_record_date'] = pd.
↳to_datetime(clean_parcel_data['parcel_record_date'], errors='coerce')

except Exception as e:
    # Show an error message if something goes wrong
    print(f"An error occurred: {str(e)}")
    import traceback
    traceback.print_exc()
```

1.2.6 Grid System, Merging, and Validation

This section performs spatial aggregation and merging of all cleaned datasets:

- **create_grid(df, x_col, y_col, cell_size):**
Splits the spatial data into uniform square grid cells (default 1km x 1km). Assigns each record a grid cell ID.
- **calc_cell(merged_df, cell_size):**
Calculates the geographic center (X, Y) of each grid cell based on its indices.
- **merge_datasets(arrests_df, lights_df, neighborhoods_df, parcels_df, cell_size):**
Combines all gridded datasets by:

- Aggregating totals (arrests, lights, parcels) per cell
- Merging neighborhood info by ward
- Calculating density per square km
- Adding spatial centroids for mapping and visualization

- **validate(df):**

Returns summary statistics like total grid cells, ward counts, and mean density values for quick data validation.

```
[13]: # Create a grid system by dividing spatial coordinates into uniform square cells
def create_grid(df, x_col, y_col, cell_size=1000):
    """Divide the area into grid cells based on X and Y coordinates."""
    df_grid = df.copy()

    # Get the full range of coordinates
    x_min, x_max = df_grid[x_col].min(), df_grid[x_col].max()
    y_min, y_max = df_grid[y_col].min(), df_grid[y_col].max()

    # Define the edges of each grid cell
    x_edges = np.arange(x_min, x_max + cell_size, cell_size)
    y_edges = np.arange(y_min, y_max + cell_size, cell_size)

    # Assign each point to a grid cell based on its X and Y position
    df_grid['grid_x_index'] = pd.cut(df_grid[x_col], bins=x_edges,
    ↪labels=range(len(x_edges) - 1))
    df_grid['grid_y_index'] = pd.cut(df_grid[y_col], bins=y_edges,
    ↪labels=range(len(y_edges) - 1))

    # Create a unique ID for each grid cell
    df_grid['grid_cell_id'] = df_grid['grid_x_index'].astype(str) + '_' +
    ↪df_grid['grid_y_index'].astype(str)

    return df_grid

# Calculate the center (X, Y) coordinates for each grid cell
def calc_cell(merged_df, cell_size=1000):
    def extract_coordinates(row):
        try:
            x_idx, y_idx = map(int, row['grid_cell_id'].split('_'))
            return pd.Series({
                'cell_center_x': x_idx * cell_size + cell_size / 2,
                'cell_center_y': y_idx * cell_size + cell_size / 2
            })
        except (ValueError, AttributeError):
            return pd.Series({'cell_center_x': np.nan, 'cell_center_y': np.nan})

    coordinates = merged_df.apply(extract_coordinates, axis=1)
    merged_df['cell_center_x'] = coordinates['cell_center_x']
```

```

merged_df['cell_center_y'] = coordinates['cell_center_y']

return merged_df

# Merge all datasets together and calculate grid-level statistics
def merge_datasets(arrests_df, lights_df, neighborhoods_df, parcels_df,
    cell_size=1000):

    print("Creating grid systems for each dataset...")
    processed_arrests = create_grid(arrests_df, 'arrest_x_cord',
    'arrest_y_cord')
    gridded_lights = create_grid(lights_df, 'light_location_x',
    'light_location_y')
    gridded_parcels = create_grid(parcels_df, 'parcel_location_x',
    'parcel_location_y')

    print("Calculating arrest statistics...")
    arrest_agg = processed_arrests.groupby('grid_cell_id').agg({
        'OBJECTID': 'count',
        'age': 'mean',
        'arrest_type': lambda x: x.mode().iloc[0] if not x.empty else None
    }).reset_index().rename(columns={
        'OBJECTID': 'total_arrests',
        'age': 'average_arrestee_age',
        'arrest_type': 'common_arrest_type'
    })

    print("Calculating light statistics...")
    light_agg = gridded_lights.groupby('grid_cell_id').agg({
        'light_device_id': 'count'
    }).reset_index().rename(columns={'light_device_id': 'total_streetlights'})

    print("Calculating parcel statistics...")
    parcel_agg = gridded_parcels.groupby('grid_cell_id').agg({
        'parcel_official_address': 'count'
    }).reset_index().rename(columns={'parcel_official_address':
    'total_parcels'})

    print("Merging aggregated statistics...")
    merged = pd.merge(arrest_agg, light_agg, on='grid_cell_id', how='outer')
    merged = pd.merge(merged, parcel_agg, on='grid_cell_id', how='outer')

    # Add ward info and neighborhood data
    distinct_ward = processed_arrests[['grid_cell_id', 'arrest_ward_number']].
    drop_duplicates()
    merged = pd.merge(merged, distinct_ward, on='grid_cell_id', how='left')

```

```

merged = pd.merge(merged, neighborhoods_df, left_on='arrest_ward_number',
↳right_on='neigh_ward_number', how='left')

# Calculate density values for each cell
cell_area_sqkm = (cell_size * cell_size) / 1_000_000
merged['arrests_per_sqkm'] = merged['total_arrests'] / cell_area_sqkm
merged['streetlights_per_sqkm'] = merged['total_streetlights'] /
↳cell_area_sqkm
merged['parcels_per_sqkm'] = merged['total_parcels'] / cell_area_sqkm

# Fill any missing data with 0
merged = merged.fillna({
    'total_arrests': 0,
    'total_streetlights': 0,
    'total_parcels': 0,
    'arrests_per_sqkm': 0,
    'streetlights_per_sqkm': 0,
    'parcels_per_sqkm': 0
})

# Add X/Y center coordinates for each grid cell
merged = calc_cell(merged, cell_size)

# Merge back the full original datasets for spatial context
final_merged = pd.merge(merged, processed_arrests, on='grid_cell_id',
↳how='left', suffixes=('', '_arrestorig'))
final_merged = pd.merge(final_merged, gridded_lights, on='grid_cell_id',
↳how='left', suffixes=('', '_lightorig'))
final_merged = pd.merge(final_merged, gridded_parcels, on='grid_cell_id',
↳how='left', suffixes=('', '_parcelorig'))

return final_merged

# Validate merged dataset by summarizing key metrics
def validate(df):
    """Check key metrics from the final dataset to ensure data is valid."""
    validation_results = {
        'total_grid_cells': len(df['grid_cell_id'].unique()),
        'unique_wards': df['arrest_ward_number'].nunique(),
        'total_arrests_sum': df['total_arrests'].sum(),
        'total_streetlights_sum': df['total_streetlights'].sum(),
        'total_parcels_sum': df['total_parcels'].sum(),
        'ward_statistics': df.groupby('arrest_ward_number').agg({
            'arrests_per_sqkm': 'mean',
            'streetlights_per_sqkm': 'mean',
            'parcels_per_sqkm': 'mean',
            'neigh_median_household_income': 'first',

```

```

        'neigh_wealth_index': 'first'
    }).head()
}
return validation_results

```

```

[14]: try:
    # Merge all datasets into one final dataset based on grid cells
    final_dataset = merge_datasets(
        clean_arrests, clean_lights, clean_neighborhoods, clean_parcel,
        ↪ cell_size=1000
    )

    # If merge was successful, display dataset structure and preview rows
    if final_dataset is not None:
        print("\nFinal Dataset Columns:")
        print(final_dataset.columns.tolist()) # List all column names

        print("\nData Sample:")
        print(final_dataset.head()) # Show the first few rows

    except NameError:

        print("Please ensure all required datasets are loaded properly.")

```

Creating grid systems for each dataset...
 Calculating arrest statistics...
 Calculating light statistics...
 Calculating parcel statistics...
 Merging aggregated statistics...

Final Dataset Columns:

```

['grid_cell_id', 'total_arrests', 'average_arrestee_age', 'common_arrest_type',
'total_streetlights', 'total_parcel', 'arrest_ward_number', 'neigh_full_name',
'neigh_ward_number', 'neigh_median_household_income',
'neigh_average_household_income', 'neigh_per_capita_income',
'neigh_total_households', 'neigh_wealth_index', 'arrests_per_sqkm',
'streetlights_per_sqkm', 'parcels_per_sqkm', 'cell_center_x', 'cell_center_y',
'OBJECTID', 'arrest_x_cord', 'arrest_y_cord', 'date_arr', 'time_arr',
'YEAR_ARR', 'age', 'race', 'arrest_zip', 'sex', 'ethnicity', 'arrest_type',
'arrest_charge_code', 'arrest_charge_description',
'arrest_ward_number_arrestorig', 'arrest_neighborhood_name',
'arrest_block_address', 'arrest_date', 'arrest_time', 'arrest_year',
'arrest_hour', 'time_period', 'grid_x_index', 'grid_y_index',
'light_location_x', 'light_location_y', 'light_address_number',
'light_street_name', 'light_operational_status', 'light_device_id',
'grid_x_index_lightorig', 'grid_y_index_lightorig', 'parcel_location_x',
'parcel_location_y', 'parcel_official_address', 'parcel_mailing_address',

```

```
'parcel_record_date', 'parcel_zipcode', 'grid_x_index_parcelorig',
'grid_y_index_parcelorig']
```

Data Sample:

	grid_cell_id	total_arrests	average_arrestee_age	common_arrest_type	\
0	0_275	0.0	NaN	NaN	
1	0_307	0.0	NaN	NaN	
2	0_83	0.0	NaN	NaN	
3	0_83	0.0	NaN	NaN	
4	0_83	0.0	NaN	NaN	

	total_streetlights	total_parcels	arrest_ward_number	neigh_full_name	\
0	0.0	0.0	NaN	NaN	
1	0.0	0.0	NaN	NaN	
2	0.0	0.0	NaN	NaN	
3	0.0	0.0	NaN	NaN	
4	0.0	0.0	NaN	NaN	

	neigh_ward_number	neigh_median_household_income	...	\
0	NaN	NaN	...	
1	NaN	NaN	...	
2	NaN	NaN	...	
3	NaN	NaN	...	
4	NaN	NaN	...	

	grid_x_index_lightorig	grid_y_index_lightorig	parcel_location_x	\
0	NaN	NaN	276311.616198	
1	NaN	NaN	276734.228534	
2	0	83	NaN	
3	0	83	NaN	
4	0	83	NaN	

	parcel_location_y	parcel_official_address	\
0	439392.248369	NaN	
1	471041.213893	NaN	
2	NaN	NaN	
3	NaN	NaN	
4	NaN	NaN	

	parcel_mailing_address	parcel_record_date	\
0	(CABEZA PRIETA NATIONAL WILDLIFE REFUGE)	NaT	
1	(CABEZA PRIETA NATIONAL WILDLIFE REFUGE)	NaT	
2	NaN	NaT	
3	NaN	NaT	
4	NaN	NaT	

	parcel_zipcode	grid_x_index_parcelorig	grid_y_index_parcelorig
0	00000	0	275

1	00000	0	307
2	NaN	NaN	NaN
3	NaN	NaN	NaN
4	NaN	NaN	NaN

[5 rows x 59 columns]

2 Data Analysis and Visualization

2.0.1 Property Crime Analysis: Night vs Day

This section analyzes property crimes based on time of day and income levels:

- Identifies property crime records using keyword matching.
- Separates crimes into night and day.
- Aggregates crime counts per grid cell.
- Calculates log-transformed rates for better analysis.
- Bins neighborhoods into income quartiles.
- Computes correlations and summary stats.
- Visualizes crime patterns using KDE plots, violin plots, regression plots, and heatmaps.

```
[15]: # Create a copy of the final dataset to work with
df = final_dataset.copy()

# Define keywords that indicate property crimes
property_crime_keywords = [
    "burglary", "theft", "larceny", "vandalism", "breaking", "robbery",
    "shoplifting", "trespass", "criminal damage", "property", "tamper",
    "organized retail", "purse snatch", "trafficking stolen property"
]

# Mark rows as property crimes if the description contains any of the keywords
df['is_property_crime'] = df['arrest_charge_description'].str.lower().str.
    ↪contains('|'.join(property_crime_keywords), na=False)

# Filter only the rows that are property crimes
df_prop = df[df['is_property_crime']].copy()

# Separate property crimes into night and day categories
df_prop_night = df_prop[df_prop['time_period'] == 'Night']
df_prop_day = df_prop[df_prop['time_period'] != 'Night']

# Count the number of property crimes per grid cell for night and day
night_crime_counts = df_prop_night.groupby('grid_cell_id').size().
    ↪reset_index(name='night_property_crime_count')
```



```

day_crime_counts = df_prop_day.groupby('grid_cell_id').size().
↳reset_index(name='day_property_crime_count')

# Merge these counts back into the main dataframe
df = df.merge(night_crime_counts, on='grid_cell_id', how='left')
df = df.merge(day_crime_counts, on='grid_cell_id', how='left')

# Replace any missing values with 0 to avoid NaNs in analysis
df['night_property_crime_count'] = df['night_property_crime_count'].fillna(0)
df['day_property_crime_count'] = df['day_property_crime_count'].fillna(0)

# Create crime-per-sqkm fields for normalization
df['night_property_crimes_per_sqkm'] = df['night_property_crime_count']
df['day_property_crimes_per_sqkm'] = df['day_property_crime_count']

# Log-transform the values to reduce skewness and handle zero values
df['log_night_property_crimes_per_sqkm'] = np.
↳log1p(df['night_property_crimes_per_sqkm'])
df['log_day_property_crimes_per_sqkm'] = np.
↳log1p(df['day_property_crimes_per_sqkm'])
df['log_streetlights_per_sqkm'] = np.log1p(df['streetlights_per_sqkm'])

# Filter dataset to only include rows with valid income data
df_income = df.dropna(subset=['neigh_median_household_income']).copy()

# Create income groups based on quartiles
df_income['income_bin'] = pd.qcut(df_income['neigh_median_household_income'],
↳q=4,
labels=['Low', 'Medium-Low', 'Medium-High',
↳'High'])

# Compute correlation matrices for day and night
corr_vars_night = ['night_property_crimes_per_sqkm', 'streetlights_per_sqkm',
↳'neigh_median_household_income']
corr_matrix_night = df[corr_vars_night].corr()

corr_vars_day = ['day_property_crimes_per_sqkm', 'streetlights_per_sqkm',
↳'neigh_median_household_income']
corr_matrix_day = df[corr_vars_day].corr()

# Display correlation matrices
print("Correlation Matrix (Day):\n", corr_matrix_day)
print("Correlation Matrix (Night):\n", corr_matrix_night)

# Get summary statistics of property crimes by income level (night and day)

```

```

income_group_stats_night = df_income.groupby('income_bin',
    ↳observed=True)['night_property_crimes_per_sqkm'].agg(['mean', 'median',
    ↳'std', 'count'])
print("Property Crime Stats by Income Level (Night):\n",
    ↳income_group_stats_night)

income_group_stats_day = df_income.groupby('income_bin',
    ↳observed=True)['day_property_crimes_per_sqkm'].agg(['mean', 'median', 'std',
    ↳'count'])
print("Property Crime Stats by Income Level (Day):\n", income_group_stats_day)

# Plotting the Data

sns.set_style('whitegrid')

# 1. KDE Plot: Compare distributions of day vs night property crimes
print("\n1. KDE distributions comparing nighttime and daytime property
    ↳crimes\n")
plt.figure(figsize=(10,6))
sns.kdeplot(df['log_night_property_crimes_per_sqkm'], fill=True,
    ↳label='Night-time Property Crimes')
sns.kdeplot(df['log_day_property_crimes_per_sqkm'], fill=True, label='Day-time
    ↳Property Crimes')
plt.title("Distribution of Log(Property Crimes per Sqkm) - Night vs. Day")
plt.xlabel("Log(Property Crimes per Sqkm)")
plt.ylabel("Density")
plt.legend()
plt.show()

# 2. Violin Plots: Show crime variation by income bin
print("\n2. Violin plots by income group for night and day property crimes
    ↳separately\n")
plt.figure(figsize=(10,6))
sns.violinplot(x='income_bin', y='log_night_property_crimes_per_sqkm',
    ↳data=df_income, inner='quartile')
plt.title("Log(Night-time Property Crimes per Sqkm) by Income Level")
plt.xlabel("Income Level")
plt.ylabel("Log(Night-time Property Crimes per Sqkm)")
plt.show()

plt.figure(figsize=(10,6))
sns.violinplot(x='income_bin', y='log_day_property_crimes_per_sqkm',
    ↳data=df_income, inner='quartile')
plt.title("Log(Day-time Property Crimes per Sqkm) by Income Level")
plt.xlabel("Income Level")

```

```

plt.ylabel("Log(Day-time Property Crimes per Sqkm)")
plt.show()

# 3. Regression Plots: Visualize streetlight-crime relationship
print("\n3. Regression plots: Streetlights vs. Property Crimes (Night vs_
    ↪Day)\n")
fig, axes = plt.subplots(1, 2, figsize=(14,6))
sns.regplot(x='log_streetlights_per_sqkm',
    ↪y='log_night_property_crimes_per_sqkm',
            data=df, scatter_kws={'alpha':0.3}, line_kws={"color":"red"},
    ↪ax=axes[0])
axes[0].set_title("Night-time Property Crimes vs. Streetlights")
axes[0].set_xlabel("Log(Streetlights per Sqkm)")
axes[0].set_ylabel("Log(Night-time Property Crimes per Sqkm)")

sns.regplot(x='log_streetlights_per_sqkm', y='log_day_property_crimes_per_sqkm',
            data=df, scatter_kws={'alpha':0.3}, line_kws={"color":"red"},
    ↪ax=axes[1])
axes[1].set_title("Day-time Property Crimes vs. Streetlights")
axes[1].set_xlabel("Log(Streetlights per Sqkm)")
axes[1].set_ylabel("Log(Day-time Property Crimes per Sqkm)")
plt.tight_layout()
plt.show()

# 4. Correlation Heatmaps: Show how variables relate for day and night
print("\n4. Heatmaps for nighttime and daytime correlations\n")
plt.figure(figsize=(12,4))
plt.subplot(1,2,1)
sns.heatmap(corr_matrix_night, annot=True, cmap='coolwarm', vmin=-1, vmax=1)
plt.title("Night-time Correlation")

plt.subplot(1,2,2)
sns.heatmap(corr_matrix_day, annot=True, cmap='coolwarm', vmin=-1, vmax=1)
plt.title("Day-time Correlation")
plt.tight_layout()
plt.show()

```

Correlation Matrix (Day):

	day_property_crimes_per_sqkm \
day_property_crimes_per_sqkm	1.000000
streetlights_per_sqkm	-0.035513
neigh_median_household_income	-0.004157

	streetlights_per_sqkm \
day_property_crimes_per_sqkm	-0.035513
streetlights_per_sqkm	1.000000
neigh_median_household_income	NaN

```


                                neigh_median_household_income
day_property_crimes_per_sqkm      -0.004157
streetlights_per_sqkm              NaN
neigh_median_household_income      1.000000
Correlation Matrix (Night):
                                night_property_crimes_per_sqkm \
night_property_crimes_per_sqkm      1.000000
streetlights_per_sqkm              -0.034470
neigh_median_household_income      0.004735

                                streetlights_per_sqkm \
night_property_crimes_per_sqkm      -0.03447
streetlights_per_sqkm              1.00000
neigh_median_household_income      NaN

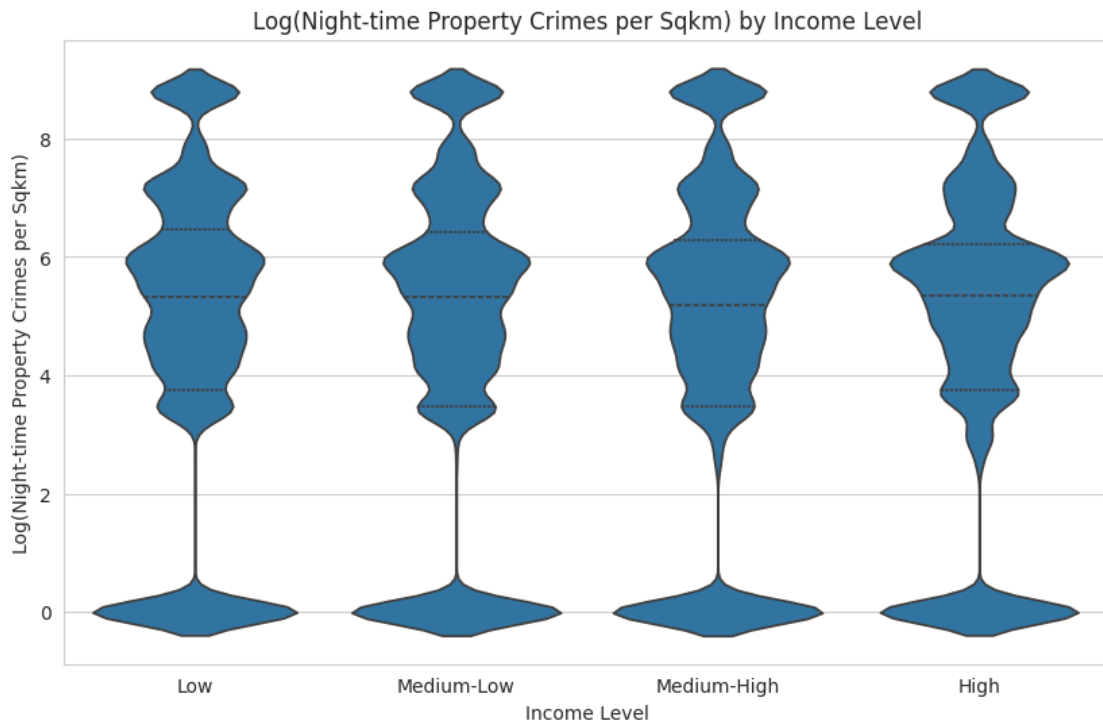
                                neigh_median_household_income
night_property_crimes_per_sqkm      0.004735
streetlights_per_sqkm              NaN
neigh_median_household_income      1.000000
Property Crime Stats by Income Level (Night):
      mean  median      std  count
income_bin
Low      866.835781   205.0  1740.061168  420774
Medium-Low  900.570956   205.0  1796.785657  425378
Medium-High  900.719244   180.0  1834.084424  412971
High      907.952683   210.0  1828.191778  414015
Property Crime Stats by Income Level (Day):
      mean  median      std  count
income_bin
Low      2906.879643   527.0  5738.318298  420774
Medium-Low  3208.164477   513.0  5992.092039  425378
Medium-High  3062.583428   472.0  6050.925476  412971
High      3020.792964   539.0  6014.906221  414015


```

1. KDE distributions comparing nighttime and daytime property crimes




2. Violin plots by income group for night and day property crimes separately

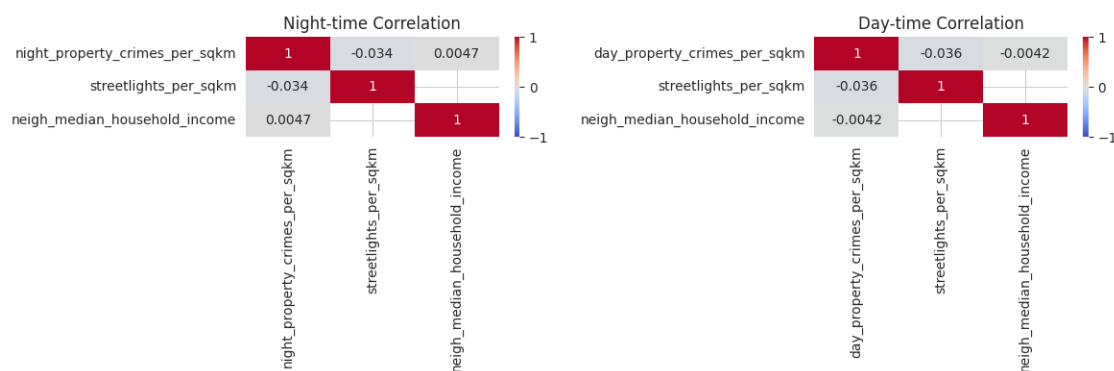




3. Regression plots: Streetlights vs. Property Crimes (Night vs Day)



4. Heatmaps for nighttime and daytime correlations



2.1 Predictive Modeling for Property Crime Analysis

This section uses machine learning to test whether streetlight density and neighborhood income are predictive of property crime rates and risk levels. We build two types of models for both daytime and nighttime:

2.1.1 1 Regression Models (Linear Regression)

Purpose:

To **predict the continuous value** of property crime rates (log-transformed per square kilometer) based on environmental and socioeconomic features.

Target Variables: - `log_night_property_crimes_per_sqkm` -
`log_day_property_crimes_per_sqkm`

Input Features: - `streetlights_per_sqkm` – Number of streetlights in a grid cell. -
`neigh_median_household_income` – Median income for the neighborhood. - `parcels_per_sqkm`
– A proxy for parcel/building density.

Model Used:

`LinearRegression()` from scikit-learn – a simple, interpretable model that assumes a linear relationship between features and the target variable.

Evaluation Metric:

Mean Squared Error (MSE) – Lower values indicate better model performance by measuring the average squared difference between predicted and actual crime rates.

2.1.2 2 Classification Models (Random Forest Classifier)

Purpose:

To **classify each grid cell** as either **High Crime Risk** or **Low Crime Risk** based on whether crime rates are above or below the median value. This helps in identifying hotspots.

Target Variables: - `crime_risk_night` (1 = above median, 0 = below median) - `crime_risk_day` (1 = above median, 0 = below median)

Input Features: Same as regression: - `streetlights_per_sqkm` - `neigh_median_household_income` - `parcels_per_sqkm`

Model Used:

`RandomForestClassifier()` – an ensemble-based classifier that combines multiple decision trees to improve accuracy and robustness. It also allows us to extract **feature importance**.

Evaluation Metric:

`classification_report()` – Shows precision, recall, f1-score, and support for both classes. Helps assess how well the model distinguishes between high and low crime areas.

2.1.3 Feature Importance Analysis

After training the Random Forest Classifiers, we compute and visualize feature importance to understand: - Which variables most strongly influence crime risk predictions. - Whether streetlight density or income level has more predictive power.

Separate plots are generated for nighttime and daytime models, making it easier to compare their behavior under different conditions.

2.1.4 Summary of Steps

1. Prepare input features and target labels for both regression and classification tasks.
2. Split data into training (80%) and testing (20%) sets.
3. Train and evaluate Linear Regression for predicting crime rates.
4. Train and evaluate Random Forest Classifier for predicting crime risk levels.
5. Visualize feature importance to interpret model behavior.

This approach gives us both **quantitative predictions** and **categorical classifications**, which are useful for planning interventions, resource allocation, and urban safety improvements.

3 Model Implementation

```
[16]: from sklearn.impute import SimpleImputer

# Copy original data
df_clean = df.copy()

# Create binary labels for high crime risk
df_clean['risk_night'] = np.
    ↳where(df_clean['log_night_property_crimes_per_sqkm'] >_
    ↳df_clean['log_night_property_crimes_per_sqkm'].median(), 1, 0)
df_clean['risk_day'] = np.where(df_clean['log_day_property_crimes_per_sqkm'] >_
    ↳df_clean['log_day_property_crimes_per_sqkm'].median(), 1, 0)
```



```

# Define features and targets
features = ['streetlights_per_sqkm', 'neigh_median_household_income',
            ↪ 'parcels_per_sqkm']
target_log_night = 'log_night_property_crimes_per_sqkm'
target_log_day = 'log_day_property_crimes_per_sqkm'
target_risk_night = 'risk_night'
target_risk_day = 'risk_day'

# Handle missing values
imputer = SimpleImputer(strategy='median')
X = imputer.fit_transform(df_clean[features])

# Targets
y_log_night = df_clean[target_log_night]
y_log_day = df_clean[target_log_day]
y_cls_night = df_clean[target_risk_night]
y_cls_day = df_clean[target_risk_day]

# Train-test splits
X_train_night, X_test_night, y_train_log_night, y_test_log_night =
    ↪ train_test_split(X, y_log_night, test_size=0.2, random_state=42)
X_train_day, X_test_day, y_train_log_day, y_test_log_day = train_test_split(X,
    ↪ y_log_day, test_size=0.2, random_state=42)
X_train_cls_night, X_test_cls_night, y_train_cls_night, y_test_cls_night =
    ↪ train_test_split(X, y_cls_night, test_size=0.2, random_state=42)
X_train_cls_day, X_test_cls_day, y_train_cls_day, y_test_cls_day =
    ↪ train_test_split(X, y_cls_day, test_size=0.2, random_state=42)

# --- Nighttime Models ---
print("### NIGHTTIME MODELS ###")

# Linear Regression
lr_night = LinearRegression()
lr_night.fit(X_train_night, y_train_log_night)
y_pred_night = lr_night.predict(X_test_night)
mse_night = mean_squared_error(y_test_log_night, y_pred_night)
print(f"Nighttime MSE (Linear Regression): {mse_night:.4f}")

# Random Forest Classifier
rf_night = RandomForestClassifier(n_estimators=100, random_state=42)
rf_night.fit(X_train_cls_night, y_train_cls_night)
y_pred_cls_night = rf_night.predict(X_test_cls_night)
print("Nighttime Crime Risk Classification Report:")
print(classification_report(y_test_cls_night, y_pred_cls_night))

# Feature importance

```

```

importance_night = pd.DataFrame({
    'Feature': features,
    'Importance': rf_night.feature_importances_
}).sort_values(by='Importance', ascending=False)

# --- Daytime Models ---
print("\n### DAYTIME MODELS ###")

# Linear Regression
lr_day = LinearRegression()
lr_day.fit(X_train_day, y_train_log_day)
y_pred_day = lr_day.predict(X_test_day)
mse_day = mean_squared_error(y_test_log_day, y_pred_day)
print(f"Daytime MSE (Linear Regression): {mse_day:.4f}")

# Random Forest Classifier
rf_day = RandomForestClassifier(n_estimators=100, random_state=42)
rf_day.fit(X_train_cls_day, y_train_cls_day)
y_pred_cls_day = rf_day.predict(X_test_cls_day)
print("Daytime Crime Risk Classification Report:")
print(classification_report(y_test_cls_day, y_pred_cls_day))

# Feature importance
importance_day = pd.DataFrame({
    'Feature': features,
    'Importance': rf_day.feature_importances_
}).sort_values(by='Importance', ascending=False)

# --- Visualization ---
plt.figure(figsize=(12, 8))

plt.subplot(2, 1, 1)
plt.barh(importance_night['Feature'], importance_night['Importance'],
         color='steelblue')
plt.title("Nighttime Crime Risk Feature Importance")
plt.xlabel("Importance")
plt.gca().invert_yaxis()

plt.subplot(2, 1, 2)
plt.barh(importance_day['Feature'], importance_day['Importance'],
         color='tomato')
plt.title("Daytime Crime Risk Feature Importance")
plt.xlabel("Importance")
plt.gca().invert_yaxis()

plt.tight_layout()
plt.show()

```

```

# Print feature importances
print("\nFeature Importances - Night:")
print(importance_night)

print("\nFeature Importances - Day:")
print(importance_day)

```

NIGHTTIME MODELS

Nighttime MSE (Linear Regression): 6.9393

Nighttime Crime Risk Classification Report:

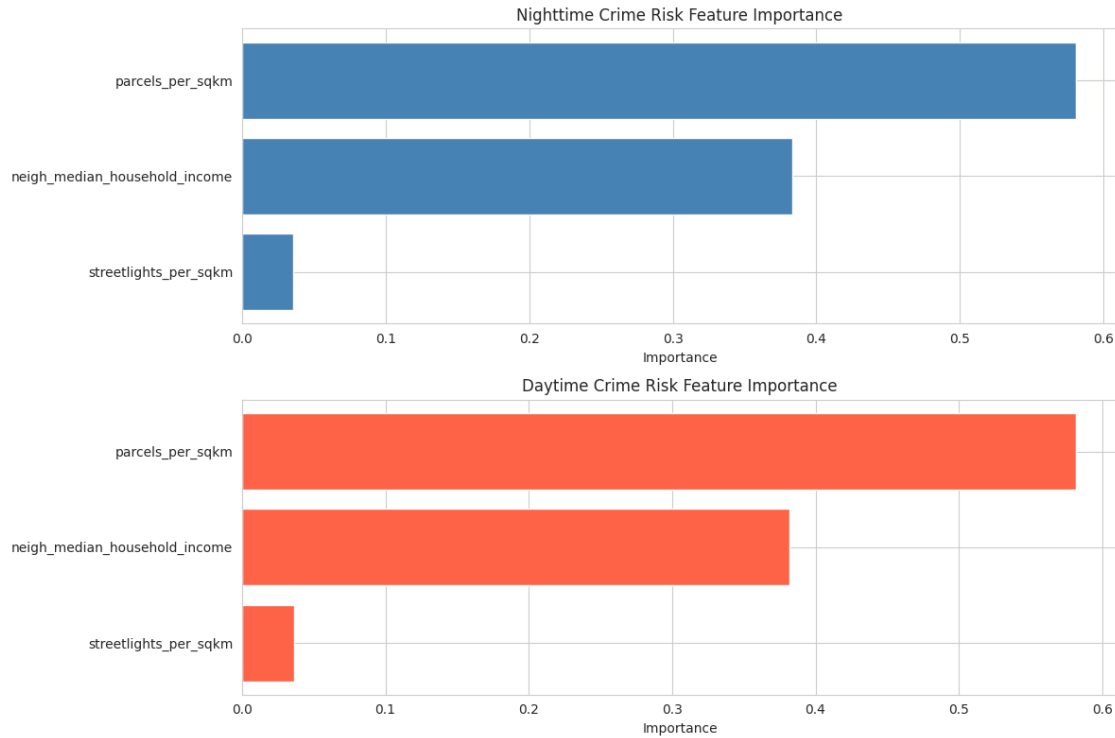
	precision	recall	f1-score	support
0	0.83	0.56	0.67	214838
1	0.67	0.88	0.76	213301
accuracy			0.72	428139
macro avg	0.75	0.72	0.71	428139
weighted avg	0.75	0.72	0.71	428139

DAYTIME MODELS

Daytime MSE (Linear Regression): 7.4918

Daytime Crime Risk Classification Report:

	precision	recall	f1-score	support
0	0.78	0.62	0.69	214654
1	0.68	0.82	0.75	213485
accuracy			0.72	428139
macro avg	0.73	0.72	0.72	428139
weighted avg	0.73	0.72	0.72	428139



Feature Importances - Night:

	Feature	Importance
2	parcels_per_sqkm	0.580722
1	neigh_median_household_income	0.383557
0	streetlights_per_sqkm	0.035722

Feature Importances - Day:

	Feature	Importance
2	parcels_per_sqkm	0.581613
1	neigh_median_household_income	0.382048
0	streetlights_per_sqkm	0.036339

4 Appropriate evaluation of the models

4.1 Appropriate Evaluation of the Models

To evaluate both regression and classification models used for predicting property crime, we use metrics that are standard and interpretable.

4.1.1 Regression Models (Linear Regression)

Metrics Used: - Mean Squared Error (MSE):

Measures the average squared difference between predicted and actual values.

→ Lower values = better model performance.

- **R² Score (Coefficient of Determination):**

Indicates how much variance in the outcome variable is explained by the model.

→ Values closer to 1 = better fit.

4.1.2 Classification Models (Random Forest Classifier)

Metrics Used: - Accuracy:

Percentage of correctly predicted labels.

- **Precision:**
Out of all predicted positives, how many were actually positive.
 - **Recall:**
Out of all actual positives, how many were correctly predicted.
 - **F1-Score:**
Harmonic mean of precision and recall.
→ Useful when dealing with imbalanced classes.
 - **Confusion Matrix:**
Visualizes how many instances were correctly and incorrectly classified as high or low risk.
-

4.1.3 Outcome of Evaluation

After this step, we can: - Understand the strengths and weaknesses of each model. - Compare day vs. night model effectiveness. - Decide if improvements (like hyperparameter tuning or more features) are needed.

```
[17]: from sklearn.metrics import r2_score, confusion_matrix, ConfusionMatrixDisplay

print("\n### NIGHTTIME MODEL EVALUATION ###\n")

# Linear Regression (Night)
r2_score_night = r2_score(y_test_log_night, y_pred_night)
print("Linear Regression (Night):")
print(f" Mean Squared Error: {mse_night:.4f}")
print(f" R2 Score: {r2_score_night:.4f}")

# Random Forest Classifier (Night)
conf_matrix_night = confusion_matrix(y_test_cls_night, y_pred_cls_night)
ConfusionMatrixDisplay(confusion_matrix=conf_matrix_night, display_labels=['Low_
↳Risk', 'High Risk']).plot(cmap='Blues')
```

```

plt.title("Confusion Matrix: Nighttime Crime Risk")
plt.show()

# -----
# DAYTIME MODEL EVALUATION
# -----
print("\n### DAYTIME MODEL EVALUATION ###\n")

# Linear Regression (Day)
r2_score_day = r2_score(y_test_log_day, y_pred_day)
print("Linear Regression (Day):")
print(f" Mean Squared Error: {mse_day:.4f}")
print(f" R2 Score: {r2_score_day:.4f}")

# Random Forest Classifier (Day)
conf_matrix_day = confusion_matrix(y_test_cls_day, y_pred_cls_day)
ConfusionMatrixDisplay(confusion_matrix=conf_matrix_day, display_labels=['Low_Risk', 'High Risk']).plot(cmap='Oranges')
plt.title("Confusion Matrix: Daytime Crime Risk")
plt.show()

# -----
# CLASSIFICATION REPORTS
# -----
print("\n### CLASSIFICATION REPORTS ###\n")

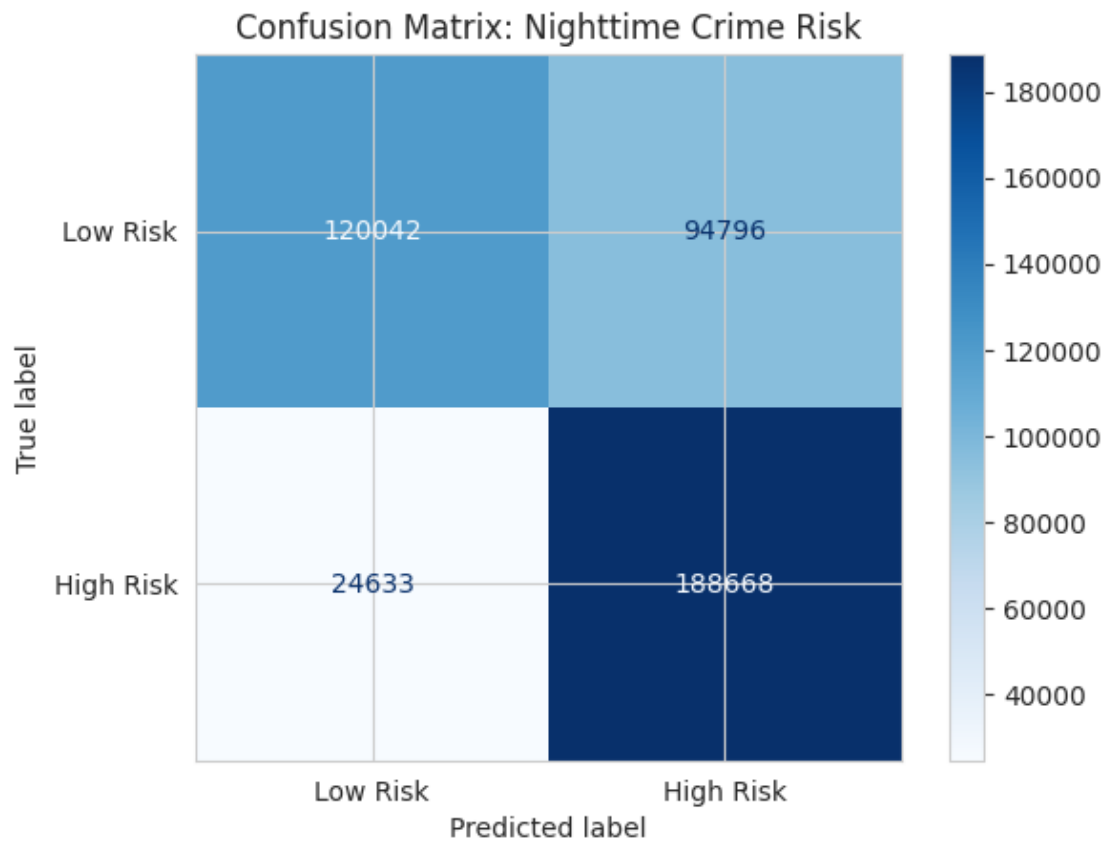
print("Random Forest Classifier - Nighttime:")
print(classification_report(y_test_cls_night, y_pred_cls_night))

print("Random Forest Classifier - Daytime:")
print(classification_report(y_test_cls_day, y_pred_cls_day))

```


NIGHTTIME MODEL EVALUATION

Linear Regression (Night):
Mean Squared Error: 6.9393
R² Score: 0.2543



DAYTIME MODEL EVALUATION

Linear Regression (Day):
Mean Squared Error: 7.4918
R² Score: 0.3356



CLASSIFICATION REPORTS

Random Forest Classifier - Nighttime:

	precision	recall	f1-score	support
0	0.83	0.56	0.67	214838
1	0.67	0.88	0.76	213301
accuracy			0.72	428139
macro avg	0.75	0.72	0.71	428139
weighted avg	0.75	0.72	0.71	428139

Random Forest Classifier - Daytime:

	precision	recall	f1-score	support
0	0.78	0.62	0.69	214654
1	0.68	0.82	0.75	213485
accuracy			0.72	428139

macro avg	0.73	0.72	0.72	428139
weighted avg	0.73	0.72	0.72	428139

4.2 Enhanced Evaluation of Crime Prediction Models

This section improves upon the basic model evaluation by incorporating advanced metrics and visual insights, allowing a deeper understanding of how the models perform during both nighttime and daytime.

4.2.1 Regression Evaluation (Night & Day)

Metrics Used: - **Mean Squared Error (MSE):** Measures average prediction error. - **R² Score:** Indicates how well the model explains variance in crime rates.

4.2.2 Classification Evaluation (Random Forest Classifier)

Standard Metrics:

- **Confusion Matrix:** Shows true positives/negatives and false positives/negatives.
- **Classification Report:** Includes precision, recall, F1-score, and support for each class.

Advanced Metrics:

- **ROC-AUC Score:** Measures model's ability to distinguish between classes.
 - A higher ROC-AUC (closer to 1) means better separation between high-risk and low-risk zones.
 - **Precision-Recall Curve:**
 - Useful when data is imbalanced.
 - Helps understand how precision and recall change at different thresholds.
-

4.2.3 Feature Importance

- Bar charts show how much each feature contributes to crime risk prediction.
 - Separate plots for nighttime and daytime improve interpretability.
-

4.2.4 Key Takeaways

- Combines multiple views (metrics + visuals) to assess model reliability.
- Helps compare day vs night predictions and identify which features are most influential.

```
[18]: from sklearn.metrics import r2_score, confusion_matrix, ConfusionMatrixDisplay, \
      ↪ roc_auc_score, precision_recall_curve
```

```

print("\n### NIGHTTIME MODEL EVALUATION ###\n")

# Linear Regression (Night)
r2_night = r2_score(y_test_log_night, y_pred_night)
print("Linear Regression (Night):")
print(f" Mean Squared Error: {mse_night:.4f}")
print(f" R2 Score: {r2_night:.4f}")

# Random Forest Classifier (Night) - Confusion Matrix
conf_matrix_night = confusion_matrix(y_test_cls_night, y_pred_cls_night)
ConfusionMatrixDisplay(confusion_matrix=conf_matrix_night, display_labels=['Low_
↳Risk', 'High Risk']).plot(cmap='Blues')
plt.title("Confusion Matrix: Nighttime Crime Risk")
plt.show()

# ROC-AUC (Night)
roc_auc_night = roc_auc_score(y_test_cls_night, rf_night.
↳predict_proba(X_test_cls_night)[: , 1])
print(f"\nRandom Forest ROC-AUC (Night): {roc_auc_night:.4f}")

# Precision-Recall Curve (Night)
precision_night, recall_night, _ = precision_recall_curve(y_test_cls_night,
↳rf_night.predict_proba(X_test_cls_night)[: , 1])
plt.figure(figsize=(8, 6))
plt.plot(recall_night, precision_night, label='Precision-Recall Curve',
↳color='blue')
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.title("Precision-Recall Curve: Nighttime")
plt.legend()
plt.show()

# -----
# DAYTIME MODEL EVALUATION
# -----
print("\n### DAYTIME MODEL EVALUATION ###\n")

# Linear Regression (Day)
r2_day = r2_score(y_test_log_day, y_pred_day)
print("Linear Regression (Day):")
print(f" Mean Squared Error: {mse_day:.4f}")
print(f" R2 Score: {r2_day:.4f}")

# Random Forest Classifier (Day) - Confusion Matrix
conf_matrix_day = confusion_matrix(y_test_cls_day, y_pred_cls_day)
ConfusionMatrixDisplay(confusion_matrix=conf_matrix_day, display_labels=['Low_
↳Risk', 'High Risk']).plot(cmap='Oranges')

```

```

plt.title("Confusion Matrix: Daytime Crime Risk")
plt.show()

# ROC-AUC (Day)
roc_auc_day = roc_auc_score(y_test_cls_day, rf_day.
    ↪predict_proba(X_test_cls_day)[: , 1])
print(f"\nRandom Forest ROC-AUC (Day): {roc_auc_day:.4f}")

# Precision-Recall Curve (Day)
precision_day, recall_day, _ = precision_recall_curve(y_test_cls_day, rf_day.
    ↪predict_proba(X_test_cls_day)[: , 1])
plt.figure(figsize=(8, 6))
plt.plot(recall_day, precision_day, label='Precision-Recall Curve',
    ↪color='darkorange')
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.title("Precision-Recall Curve: Daytime")
plt.legend()
plt.show()

# -----
# CLASSIFICATION REPORTS
# -----
print("\n### CLASSIFICATION REPORTS ###\n")

print("Random Forest Classifier - Nighttime:")
print(classification_report(y_test_cls_night, y_pred_cls_night))

print("Random Forest Classifier - Daytime:")
print(classification_report(y_test_cls_day, y_pred_cls_day))

# -----
# FEATURE IMPORTANCE VISUALIZATION
# -----
plt.figure(figsize=(14, 10))

# Nighttime Feature Importance
plt.subplot(2, 1, 1)
plt.barh(importance_night['Feature'], importance_night['Importance'],
    ↪color='steelblue')
plt.xlabel("Importance")
plt.title("Feature Importance - Nighttime Crime Risk")
plt.gca().invert_yaxis()

# Daytime Feature Importance
plt.subplot(2, 1, 2)

```

```
plt.barh(importance_day['Feature'], importance_day['Importance'],
         color='tomato')
plt.xlabel("Importance")
plt.title("Feature Importance - Daytime Crime Risk")
plt.gca().invert_yaxis()

plt.tight_layout()
plt.show()

# Print feature importance tables
print("\nFeature Importance - Nighttime:")
print(importance_night)


print("\nFeature Importance - Daytime:")
print(importance_day)
```

NIGHTTIME MODEL EVALUATION


Linear Regression (Night):

Mean Squared Error: 6.9393

R² Score: 0.2543



Random Forest ROC-AUC (Night): 0.7694




DAYTIME MODEL EVALUATION

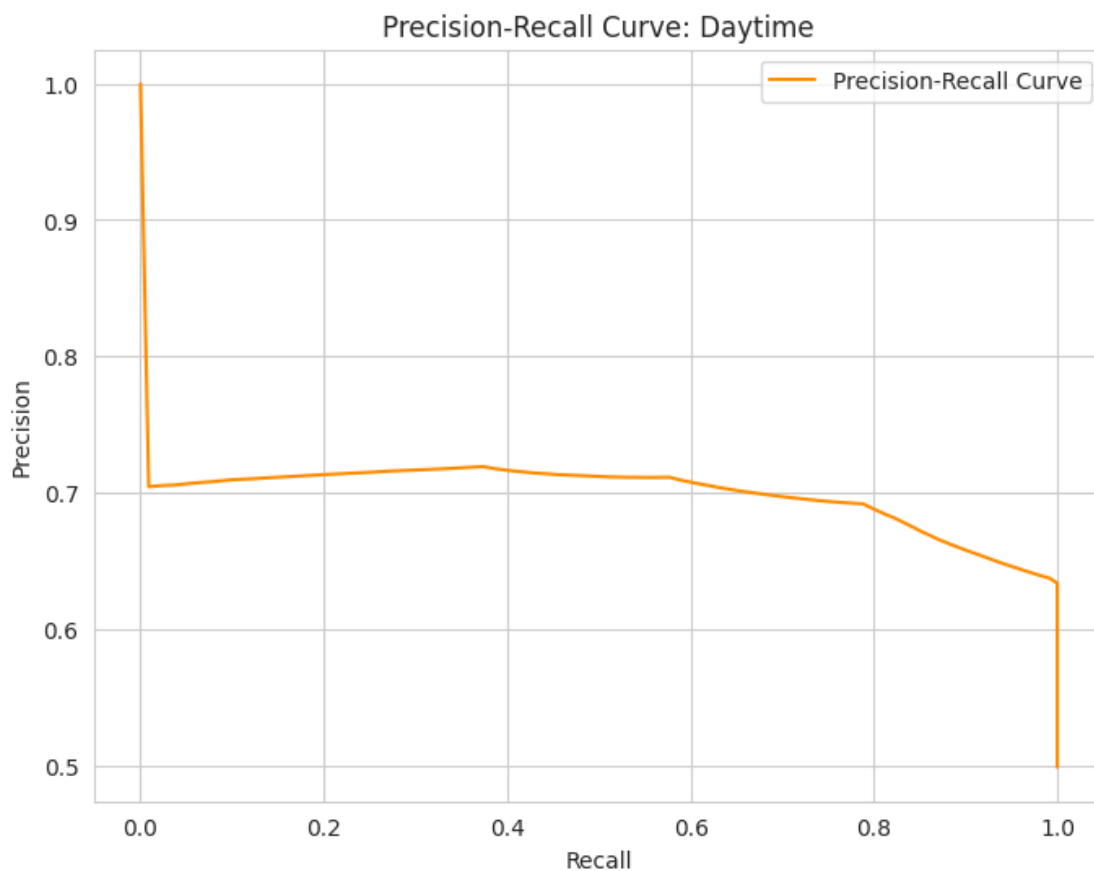
Linear Regression (Day):

Mean Squared Error: 7.4918

R^2 Score: 0.3356



Random Forest ROC-AUC (Day): 0.7746



CLASSIFICATION REPORTS


Random Forest Classifier - Nighttime:

	precision	recall	f1-score	support
0	0.83	0.56	0.67	214838
1	0.67	0.88	0.76	213301
accuracy			0.72	428139
macro avg	0.75	0.72	0.71	428139
weighted avg	0.75	0.72	0.71	428139

Random Forest Classifier - Daytime:

	precision	recall	f1-score	support
0	0.78	0.62	0.69	214654
1	0.68	0.82	0.75	213485
accuracy			0.72	428139

macro avg	0.73	0.72	0.72	428139
weighted avg	0.73	0.72	0.72	428139



Feature Importance - Nighttime:

	Feature	Importance
2	parcels_per_sqkm	0.580722
1	neigh_median_household_income	0.383557
0	streetlights_per_sqkm	0.035722

Feature Importance - Daytime:

	Feature	Importance
2	parcels_per_sqkm	0.581613
1	neigh_median_household_income	0.382048
0	streetlights_per_sqkm	0.036339

4.3 Analyzing the Correlation Between Median Household Income and Neighborhood Crimerate in Tucson

Problem Statement

Poverty, homelessness, and the divide between the top and bottom earners is a pressing issue in some parts of Tucson. There are many country wide statistics examining the correlation between income and crime, but not specifically in Tucson.

Hypothesis

Neighborhoods with lower median household income will have a higher number of crimes than higher income areas.

Objective

This section aims to:

- Determine if a measurable correlation exists between median household income and crime rates in Tucson.
 - Use machine learning models to predict crime risk in low income areas to help address areas that need the most attention.
-

4.3.1 Data Gathering

Overview

Data for this project was sourced from public governmental Tucson datasets in CSV format and includes: - Arrest records (Tucson Police Department) in 2021 - Neighborhood income data

```
[19]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
from IPython.display import display

#load data frames from csv
# neighborhoodPath = "Neighborhood_Income.csv"
df_neighborhood_income = tuc_neighbourhood_income
df_neighborhood_income = df_neighborhood_income[df_neighborhood_income["HasData_1"] != 0]

# crimePath = "Tucson_Police_Arrests_-_2021_-_Open_Data.csv"
df_crime = tuc_arrest_data

print("crime head:")
display(df_crime.head())

print("neighborhood head:")
display(df_neighborhood_income.head())
```

crime head:

	OBJECTID	X	Y	arre_id	case_id	agency	\
0	1	9.900089e+05	470751.276735	2021000107	2101020104	TPD	
1	2	9.900089e+05	470751.276735	2021000107	2101020104	TPD	
2	3	9.900089e+05	470751.276735	2021000107	2101020104	TPD	

3	4	1.053154e+06	443419.380064	2021000110	2101020138	TPD
4	5	1.053154e+06	443419.380064	2021000110	2101020138	TPD

		date_arr	time_arr	datetime_arr	MONTH_ARR	...	\
0	2021/01/02	00:00:00+00	1731	2021/01/02 17:31:00+00	01-Jan	...	
1	2021/01/02	00:00:00+00	1731	2021/01/02 17:31:00+00	01-Jan	...	
2	2021/01/02	00:00:00+00	1731	2021/01/02 17:31:00+00	01-Jan	...	
3	2021/01/02	00:00:00+00	1844	2021/01/02 18:44:00+00	01-Jan	...	
4	2021/01/02	00:00:00+00	1844	2021/01/02 18:44:00+00	01-Jan	...	

	LOC_STATUS	WARD	NHA_NAME	TMSECT	DIVISION	DIVISION_NO	\
0	GEOCODED	3.0	NaN	NaN	Operations Division West	T2	
1	GEOCODED	3.0	NaN	NaN	Operations Division West	T2	
2	GEOCODED	3.0	NaN	NaN	Operations Division West	T2	
3	GEOCODED	2.0	Eastside	NaN	Operations Division East	T4	
4	GEOCODED	2.0	Eastside	NaN	Operations Division East	T4	

	DIVSECT	TRSQ	City_geo	ADDRESS_100BLK
0	T203	13S13E24NW	TUCSON	4598 N ORACLE RD
1	T203	13S13E24NW	TUCSON	4598 N ORACLE RD
2	T203	13S13E24NW	TUCSON	4598 N ORACLE RD
3	T406	14S15E14NE	TUCSON	10198 E ESSEX VILLAGE DR
4	T406	14S15E14NE	TUCSON	10198 E ESSEX VILLAGE DR

[5 rows x 39 columns]

neighborhood head:

	OBJECTID	NAME	WARD	DATASOURCE	ID	sourceCountry	\
0	1	A Mountain	1	NEIGHBORHOODS	0	US	
1	2	Adelanto	3	NEIGHBORHOODS	1	US	
2	3	Alvernon Heights	5	NEIGHBORHOODS	2	US	
3	4	Amphi	3	NEIGHBORHOODS	3	US	
4	5	Armory Park	6	NEIGHBORHOODS	4	US	

	ENRICH_FID	aggregationMethod	\
0	1	BlockApportionment:US.BlockGroups	
1	2	BlockApportionment:US.BlockGroups	
2	3	BlockApportionment:US.BlockGroups	
3	4	BlockApportionment:US.BlockGroups	
4	5	BlockApportionment:US.BlockGroups	

	populationToPolygonSizeRating	apportionmentConfidence	...	AGGDIA75CY	\
0	2.191	2.576	...	1590160	
1	2.191	2.576	...	154598	
2	2.191	2.576	...	172634	
3	2.191	2.576	...	2760918	
4	2.191	2.576	...	3785750	

	ID_1	sourceCountry_1	ENRICH_FID_1	aggregationMethod_1	\
0	0	US	1	BlockApportionment:US.BlockGroups	
1	1	US	2	BlockApportionment:US.BlockGroups	
2	2	US	3	BlockApportionment:US.BlockGroups	
3	3	US	4	BlockApportionment:US.BlockGroups	
4	4	US	5	BlockApportionment:US.BlockGroups	

	populationToPolygonSizeRating_1	apportionmentConfidence_1	HasData_1	\
0		2.191	2.576	1
1		2.191	2.576	1
2		2.191	2.576	1
3		2.191	2.576	1
4		2.191	2.576	1

	TOTHH_CY	WLTHINDXCY
0	1103	32
1	117	28
2	99	26
3	3105	20
4	1223	48

[5 rows x 168 columns]

4.3.2 Visualisations

- **Plot the number of crimes per neighborhood:**

First drop any duplicates in the data and plot unique arrests per neighborhood.

Only show the first and last 20 neighborhoods.

```
[20]: # Bar chart of crimes per neighborhood
df_crime.drop_duplicates("arre_id", inplace=True)
crimes_per_neighborhood = df_crime[["NHA_NAME"]].value_counts()

topcrimes_per_neighborhood = crimes_per_neighborhood.head(20)[::-1]
botcrimes_per_neighborhood = crimes_per_neighborhood.tail(20)[::-1]


neighborhoodCrimes = [n[0] for n in crimes_per_neighborhood.index]

topneighborhoodCrimes = neighborhoodCrimes[:20][::-1]
botneighborhoodCrimes = neighborhoodCrimes[-20:][::-1]

plt.figure(figsize=(10, 22))
plt.barh(botneighborhoodCrimes, botcrimes_per_neighborhood, color = "b")
plt.barh(topneighborhoodCrimes, topcrimes_per_neighborhood, color = "r")
```

```
plt.legend(["Bottom 20", "Top 20"])
plt.ylabel("Neighborhood")
plt.xlabel("Reported Unique Arrests")
plt.title("Reported Arrests per Neighborhood in Tucson, AZ")
plt.grid(axis="x")
plt.show()
```

Reported Arrests per Neighborhood in Tucson, AZ



- **Plot the median income per neighborhood:**

For comparison this also shows the top and bottom 20 neighborhoods.

You may notice some overlap with the top and bottom and vice versa.


```
[21]: #Get median household income from dataframe and set the index to the name from
↳2010 - 2019
df_neighborhood_income_name = df_neighborhood_income[["NAME", "MEDHINC_CY"]]
df_neighborhood_income_name.set_index("NAME", inplace=True)
sorted_df_top_20 = df_neighborhood_income_name.sort_values("MEDHINC_CY",
↳ascending=True).head(20)
sorted_df_bottom_20 = df_neighborhood_income_name.sort_values("MEDHINC_CY",
↳ascending=True).tail(20)

#source -> https://www.census.gov/quickfacts/fact/table/tucsoncityarizona/
↳INC110223 => 2019 - 2023
print(f"Median Household Income of all of Tuscon's neighborhood: $54,546")

plt.figure(figsize=(10, 22))
plt.xticks(rotation=45)
plt.barh(sorted_df_top_20.index, sorted_df_top_20["MEDHINC_CY"], color="r")
plt.barh(sorted_df_bottom_20.index, sorted_df_bottom_20["MEDHINC_CY"],
↳color="b")
plt.xlabel("Neighborhood Income")
plt.ylabel("Neighborhood")
plt.title("Top and Bottom 20 Neighborhoods by Median Household Income")
plt.grid(axis="x")
plt.legend(["Bottom 20", "Top 20"])
plt.show()
```

Median Household Income of all of Tuscon's neighborhood: \$54,546

Top and Bottom 20 Neighborhoods by Median Household Income



- **Data cleaning step:**

Make dataframes based on necessary information and besides dropping duplicates, remove empty or irrelevant data.

```
[22]: #Data cleaning for correlation between neighborhood household income and
      ↪neighborhood crime rate
crimes_per_neighborhood_df = df_crime[["NHA_NAME"]].value_counts().
      ↪to_frame(name="CRIMES")
crimes_per_neighborhood_df.index.rename("NAME", inplace=True)

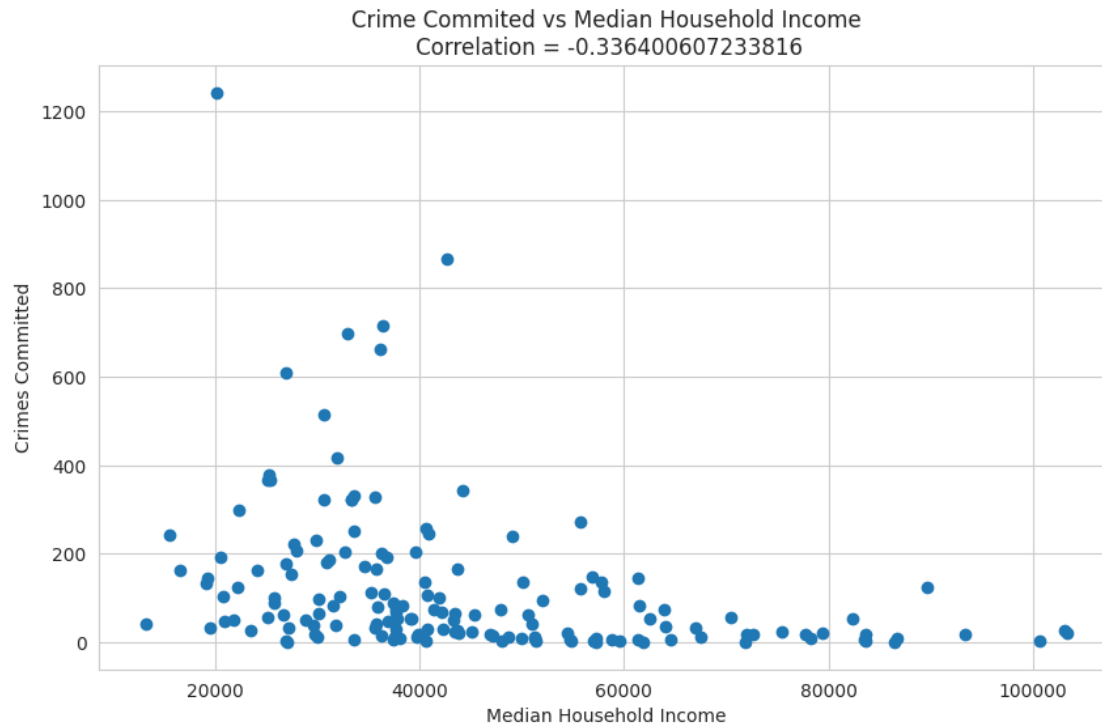
crime_and_income_df = crimes_per_neighborhood_df.
      ↪merge(df_neighborhood_income_name, on="NAME", how="outer")
crime_and_income_df.sort_index(axis=0, ascending=True, inplace=True)
#dropped 3 neighborhoods because there was NaN data for a column, ie either
      ↪crime/ income wasn't present
crime_and_income_df = crime_and_income_df.dropna()
```

4.3.3 Crime Analysis: Income vs Crime

This section provides visualisations to show the correlation between the number of crimes and median household income of Tucson Neighborhoods:

```
[23]: # Scatter plot for correlation between income and crime rate
correlation = crime_and_income_df["MEDHINC_CY"].
      ↪corr(crime_and_income_df["CRIMES"])

plt.figure(figsize=(10, 6))
plt.scatter(crime_and_income_df["MEDHINC_CY"], crime_and_income_df["CRIMES"] )
plt.xlabel("Median Household Income")
plt.ylabel("Crimes Committed")
plt.title(f"Crime Committed vs Median Household Income\n Correlation =
      ↪{correlation}")
plt.show()
```

4.3.4 Regression Model (Linear Regression)

We will try a linear regression model as it seems to follow a decreasing trend.

20% of the data will be used as testing data.

Metrics Used: - Root Mean Squared Error (RMSE):

Measures the average difference between predicted and actual values.

→ Lower values = better model performance.

- **R² Score (Coefficient of Determination):**

Indicates how much variance in the outcome variable is explained by the model.

→ Values closer to 1 = better fit.

```
[24]: # first model will be a linear regression model
from sklearn.model_selection import train_test_split
from sklearn import linear_model
from sklearn import metrics
from sklearn.metrics import mean_squared_error, r2_score

x = crime_and_income_df[["MEDHINC_CY"]]
y = crime_and_income_df["CRIMES"]

x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2)
```

```

regr = linear_model.LinearRegression()

regr.fit(x_train, y_train)

y_pred = regr.predict(x_test)


linear_rmse = np.sqrt(mean_squared_error(y_test, y_pred))
linear_r2 = r2_score(y_test, y_pred)

# plot the data
plt.scatter(x_test, y_test, color='black')
plt.plot(x_test, y_pred, color = 'b', linewidth=3)

plt.title("Linear Regression on Neighborhood Income vs Crimerate")
plt.xlabel("Median Household Neighborhood Income")
plt.ylabel("# of Crimes Committed")

plt.xticks()
plt.yticks()
plt.show()

```



4.3.5 Regression Model (Logarithmic Regression)

It appears we can do better than a linear regression so let's try logarithmic regression to fit the logarithmic curve the data has.

20% of the data will be used as testing data.

Same Metrics Used to Compare: - Root Mean Squared Error (RMSE):

Measures the average difference between predicted and actual values.

→ Lower values = better model performance.

- **R² Score (Coefficient of Determination):**

Indicates how much variance in the outcome variable is explained by the model.

→ Values closer to 1 = better fit.

```
[25]: # This seems to not be the best fit because of outliers in the lower income_
      ↪range
      # We can try a logarithmic regression model to better suit it.

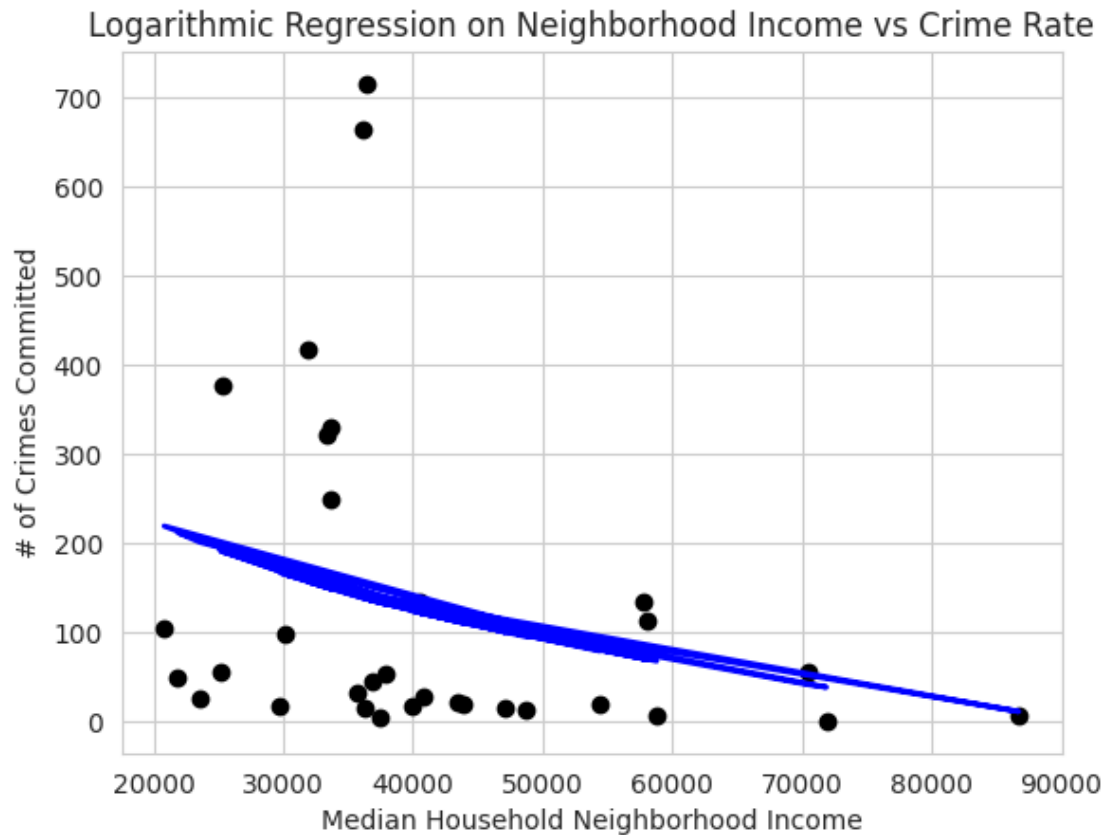
x_train, x_test, y_train, y_test = train_test_split(np.log(x), y, test_size=0.2)

regr.fit(x_train, y_train)

y_pred = regr.predict(x_test)

log_rmse = np.sqrt(mean_squared_error(y_test, y_pred))
log_r2 = r2_score(y_test, y_pred)

plt.title("Logarithmic Regression on Neighborhood Income vs Crime Rate")
plt.scatter(np.exp(x_test), y_test, color='black')
plt.plot(np.exp(x_test), y_pred, color='b', linewidth=2)
plt.xlabel("Median Household Neighborhood Income")
plt.ylabel("# of Crimes Committed")
plt.xticks()
plt.yticks()
plt.show()
```



Compare metrics for the linear vs logarithmic regression

Metrics Used:

- Root Mean Squared Error (RMSE)
- R^2 Score

These results show that logarithmic regression is much better for this data.

- As stated earlier, a lower MSE (and RMSE which is just taking the square root of that to be more accurate to the data) signifies that it is a better model for the data.
- Also, the R^2 is higher showing that it fits the variance in data better. It is still not very high because of the outliers in the lower income neighborhoods.

```
[26]: print("linear regression:")
      print(f"RMSE = {linear_rmse}")
      print(f"R^2 = {linear_r2}")
      print("\nlogarithmic regression:")
      print(f"RMSE = {log_rmse}")
      print(f"R^2 = {log_r2}")
```

```
linear regression:
RMSE = 166.81496237791615
R^2 = 0.049218622096143916
```

```
logarithmic regression:
RMSE = 178.54820078463678
R^2 = 0.06085131171673752
```

```
[30]: from google.colab import drive
drive.mount('/content/drive')
!ls "/content/drive/MyDrive/Colab_Notebook/"
!apt-get install pandoc

!apt-get install -y texlive-xetex texlive-fonts-recommended texlive-latex-extra
!pip install nbconvert
!jupyter nbconvert --to pdf "/content/drive/MyDrive/Colab Notebooks/Final_
↳Project csc 380.ipynb" --output "/content/Final_Project_csc_380.pdf"

from google.colab import files
files.download("/content/Final_Project_csc_380.pdf")
```

```
Drive already mounted at /content/drive; to attempt to forcibly remount, call
drive.mount("/content/drive", force_remount=True).
ls: cannot access '/content/drive/MyDrive/Colab_Notebook/': No such file or
directory
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
texlive-fonts-recommended is already the newest version (2021.20220204-1).
texlive-latex-extra is already the newest version (2021.20220204-1).
texlive-xetex is already the newest version (2021.20220204-1).
0 upgraded, 0 newly installed, 0 to remove and 34 not upgraded.
Requirement already satisfied: nbconvert in /usr/local/lib/python3.11/dist-
packages (7.16.6)
Requirement already satisfied: beautifulsoup4 in /usr/local/lib/python3.11/dist-
packages (from nbconvert) (4.13.4)
Requirement already satisfied: bleach!=5.0.0 in /usr/local/lib/python3.11/dist-
packages (from bleach[css]!=5.0.0->nbconvert) (6.2.0)
Requirement already satisfied: defusedxml in /usr/local/lib/python3.11/dist-
packages (from nbconvert) (0.7.1)
Requirement already satisfied: jinja2>=3.0 in /usr/local/lib/python3.11/dist-
packages (from nbconvert) (3.1.6)
Requirement already satisfied: jupyter-core>=4.7 in
/usr/local/lib/python3.11/dist-packages (from nbconvert) (5.7.2)
Requirement already satisfied: jupyterlab-pygments in
/usr/local/lib/python3.11/dist-packages (from nbconvert) (0.3.0)
Requirement already satisfied: markupsafe>=2.0 in
```

/usr/local/lib/python3.11/dist-packages (from nbconvert) (3.0.2)
 Requirement already satisfied: mistune<4,>=2.0.3 in
 /usr/local/lib/python3.11/dist-packages (from nbconvert) (3.1.3)
 Requirement already satisfied: nbclient>=0.5.0 in
 /usr/local/lib/python3.11/dist-packages (from nbconvert) (0.10.2)
 Requirement already satisfied: nbformat>=5.7 in /usr/local/lib/python3.11/dist-
 packages (from nbconvert) (5.10.4)
 Requirement already satisfied: packaging in /usr/local/lib/python3.11/dist-
 packages (from nbconvert) (24.2)
 Requirement already satisfied: pandocfilters>=1.4.1 in
 /usr/local/lib/python3.11/dist-packages (from nbconvert) (1.5.1)
 Requirement already satisfied: pygments>=2.4.1 in
 /usr/local/lib/python3.11/dist-packages (from nbconvert) (2.19.1)
 Requirement already satisfied: traitlets>=5.1 in /usr/local/lib/python3.11/dist-
 packages (from nbconvert) (5.7.1)
 Requirement already satisfied: webencodings in /usr/local/lib/python3.11/dist-
 packages (from bleach!=5.0.0->bleach[css]!=5.0.0->nbconvert) (0.5.1)
 Requirement already satisfied: tinycss2<1.5,>=1.1.0 in
 /usr/local/lib/python3.11/dist-packages (from bleach[css]!=5.0.0->nbconvert)
 (1.4.0)
 Requirement already satisfied: platformdirs>=2.5 in
 /usr/local/lib/python3.11/dist-packages (from jupyter-core>=4.7->nbconvert)
 (4.3.7)
 Requirement already satisfied: jupyter-client>=6.1.12 in
 /usr/local/lib/python3.11/dist-packages (from nbclient>=0.5.0->nbconvert)
 (6.1.12)
 Requirement already satisfied: fastjsonschema>=2.15 in
 /usr/local/lib/python3.11/dist-packages (from nbformat>=5.7->nbconvert) (2.21.1)
 Requirement already satisfied: jsonschema>=2.6 in
 /usr/local/lib/python3.11/dist-packages (from nbformat>=5.7->nbconvert) (4.23.0)
 Requirement already satisfied: soupsieve>1.2 in /usr/local/lib/python3.11/dist-
 packages (from beautifulsoup4->nbconvert) (2.7)
 Requirement already satisfied: typing-extensions>=4.0.0 in
 /usr/local/lib/python3.11/dist-packages (from beautifulsoup4->nbconvert)
 (4.13.2)
 Requirement already satisfied: attrs>=22.2.0 in /usr/local/lib/python3.11/dist-
 packages (from jsonschema>=2.6->nbformat>=5.7->nbconvert) (25.3.0)
 Requirement already satisfied: jsonschema-specifications>=2023.03.6 in
 /usr/local/lib/python3.11/dist-packages (from
 jsonschema>=2.6->nbformat>=5.7->nbconvert) (2025.4.1)
 Requirement already satisfied: referencing>=0.28.4 in
 /usr/local/lib/python3.11/dist-packages (from
 jsonschema>=2.6->nbformat>=5.7->nbconvert) (0.36.2)
 Requirement already satisfied: rpds-py>=0.7.1 in /usr/local/lib/python3.11/dist-
 packages (from jsonschema>=2.6->nbformat>=5.7->nbconvert) (0.24.0)
 Requirement already satisfied: pyzmq>=13 in /usr/local/lib/python3.11/dist-
 packages (from jupyter-client>=6.1.12->nbclient>=0.5.0->nbconvert) (24.0.1)
 Requirement already satisfied: python-dateutil>=2.1 in

```

/usr/local/lib/python3.11/dist-packages (from jupyter-
client>=6.1.12->nbclient>=0.5.0->nbconvert) (2.9.0.post0)
Requirement already satisfied: tornado>=4.1 in /usr/local/lib/python3.11/dist-
packages (from jupyter-client>=6.1.12->nbclient>=0.5.0->nbconvert) (6.4.2)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.11/dist-
packages (from python-dateutil>=2.1->jupyter-
client>=6.1.12->nbclient>=0.5.0->nbconvert) (1.17.0)
[NbConvertApp] Converting notebook /content/drive/MyDrive/Colab Notebooks/Final
Project csc 380.ipynb to pdf
[NbConvertApp] ERROR | Error while converting '/content/drive/MyDrive/Colab
Notebooks/Final Project csc 380.ipynb'
Traceback (most recent call last):
  File "/usr/local/lib/python3.11/dist-packages/nbconvert/nbconvertapp.py", line
487, in export_single_notebook
    output, resources = self.exporter.from_filename(
                        ~~~~~~

  File "/usr/local/lib/python3.11/dist-
packages/nbconvert/exporters/templateexporter.py", line 390, in from_filename
    return super().from_filename(filename, resources, **kw) #
type:ignore[return-value]
    ~~~~~~

  File "/usr/local/lib/python3.11/dist-
packages/nbconvert/exporters/exporter.py", line 201, in from_filename
    return self.from_file(f, resources=resources, **kw)
    ~~~~~~

  File "/usr/local/lib/python3.11/dist-
packages/nbconvert/exporters/templateexporter.py", line 396, in from_file
    return super().from_file(file_stream, resources, **kw) #
type:ignore[return-value]
    ~~~~~~

  File "/usr/local/lib/python3.11/dist-
packages/nbconvert/exporters/exporter.py", line 220, in from_file
    return self.from_notebook_node(
    ~~~~~~

  File "/usr/local/lib/python3.11/dist-packages/nbconvert/exporters/pdf.py",
line 184, in from_notebook_node
    latex, resources = super().from_notebook_node(nb, resources=resources, **kw)
    ~~~~~~

  File "/usr/local/lib/python3.11/dist-packages/nbconvert/exporters/latex.py",
line 92, in from_notebook_node
    return super().from_notebook_node(nb, resources, **kw)
    ~~~~~~

  File "/usr/local/lib/python3.11/dist-
packages/nbconvert/exporters/templateexporter.py", line 429, in
from_notebook_node
    output = self.template.render(nb=nb_copy, resources=resources)
    ~~~~~~

  File "/usr/local/lib/python3.11/dist-packages/jinja2/environment.py", line

```

```

1295, in render
    self.environment.handle_exception()
  File "/usr/local/lib/python3.11/dist-packages/jinja2/environment.py", line
942, in handle_exception
    raise rewrite_traceback_stack(source=source)
  File "/usr/local/share/jupyter/nbconvert/templates/latex/index.tex.j2", line
8, in top-level template code
    ((* extends cell_style *))
    ~~~~~

  File
"/usr/local/share/jupyter/nbconvert/templates/latex/style_jupyter.tex.j2", line
176, in top-level template code
    \prompt{(((prompt)))}{(((prompt_color)))}{(((execution_count)))}{(((extra_sp
ace)))}
    ~~~~~

  File "/usr/local/share/jupyter/nbconvert/templates/latex/base.tex.j2", line 7,
in top-level template code
    ((* extends 'document_contents.tex.j2' -*))
    ~~~~~

  File
"/usr/local/share/jupyter/nbconvert/templates/latex/document_contents.tex.j2",
line 51, in top-level template code
    ((* block figure scoped -*))
    ~~~~~

  File "/usr/local/share/jupyter/nbconvert/templates/latex/display_priority.j2",
line 5, in top-level template code
    ((* extends 'null.j2' -*))
    ~~~~~

  File "/usr/local/share/jupyter/nbconvert/templates/latex/null.j2", line 30, in
top-level template code
    ((* block body -*))
  File "/usr/local/share/jupyter/nbconvert/templates/latex/base.tex.j2", line
241, in block 'body'
    ((( super() )))
  File "/usr/local/share/jupyter/nbconvert/templates/latex/null.j2", line 32, in
block 'body'
    ((* block any_cell scoped -*))
    ~~~~~

  File "/usr/local/share/jupyter/nbconvert/templates/latex/null.j2", line 85, in
block 'any_cell'
    ((* block markdowncell scoped-*)) ((* endblock markdowncell -*))
    ~~~~~

  File
"/usr/local/share/jupyter/nbconvert/templates/latex/document_contents.tex.j2",
line 68, in block 'markdowncell'
    ((( cell.source | citation2latex | strip_files_prefix |
convert_pandoc('markdown+tex_math_double_backslash', 'json',extra_args=[]) |
resolve_references | convert_explicitly_relative_paths |

```



```

convert_pandoc('json','latex'))))
~~~~~
File "/usr/local/lib/python3.11/dist-packages/nbconvert/filters/pandoc.py",
line 36, in convert_pandoc
    return pandoc(source, from_format, to_format, extra_args=extra_args)
~~~~~
File "/usr/local/lib/python3.11/dist-packages/nbconvert/utils/pandoc.py", line
50, in pandoc
    check_pandoc_version()
File "/usr/local/lib/python3.11/dist-packages/nbconvert/utils/pandoc.py", line
98, in check_pandoc_version
    v = get_pandoc_version()
~~~~~
File "/usr/local/lib/python3.11/dist-packages/nbconvert/utils/pandoc.py", line
75, in get_pandoc_version
    raise PandocMissing()
nbconvert.utils.pandoc.PandocMissing: Pandoc wasn't found.
Please check that pandoc is installed:
https://pandoc.org/installing.html

```

```

-----
FileNotFoundError                                Traceback (most recent call last)
<ipython-input-30-271b39cf7e26> in <cell line: 0>()
      9
     10 from google.colab import files
--> 11 files.download("/content/Final_Project_csc_380.pdf")
     12

/usr/local/lib/python3.11/dist-packages/google/colab/files.py in
↳ download(filename)
     231     if not _os.path.exists(filename):
     232         msg = 'Cannot find file: {}'.format(filename)
--> 233         raise FileNotFoundError(msg) # pylint: disable=undefined-variable
     234
     235     comm_manager = _IPython.get_ipython().kernel.comm_manager

FileNotFoundError: Cannot find file: /content/Final_Project_csc_380.pdf

```